

MPRI FUN
Towards
machine-
checked
proofs

François
Pottier

Mechanization

Rocq in a
nutshell

Syntax with
binders

de Bruijn
Nominal

Towards machine-checked proofs

MPRI FUN

François Pottier



2025–2026

Why formalize programming languages?

To obtain **precise definitions** of programming languages.

To obtain **rigorous proofs** of soundness for tools such as

- interpreters,
- compilers,
- type systems (“well-typed programs do not go wrong”),
- type-checkers and type inference engines,
- static analyzers (e.g. abstract interpreters),
- program logics (e.g. Hoare logic, separation logic),
- deductive program provers (e.g. verification condition generators).

Challenge 1: Scale

Hand-written proofs have difficulty **scaling up**:

- From minimal calculi (λ , π) and toy languages (IMP, MiniML) to large **real-world languages** such as Java, C, JavaScript, ...
- From textbook compilers to multi-pass **optimizing compilers** producing code for real processors.
- From textbook abstract interpreters to **scalable and precise static analyzers** such as Astrée.

Challenge 2: Trust

Hand-written proofs are seldom **trustworthy**.

- Authors **struggle** with huge LaTeX documents.
- Reviewers **give up** on checking huge but boring proofs.
- Proof cases are **omitted** because they are “obvious”.
- It is difficult to **maintain** hand-written proofs as definitions evolve.

Opportunity: machine-assisted proof

Mechanized theorem proving has made great progress.

Landmark examples in mathematics:

- the 4-color th.: Haken & Appel (1976), Gonthier & Werner (2005);
- the Feit-Thompson theorem: Gonthier *et al.* (2013);
- Kepler's conjecture: Hales *et al.* (2015);
- The Liquid Tensor Experiment: Commelin *et al.* (2022).

Programming language theory is a [good match](#) for proof assistants:

- discrete objects (trees); no reals, no analysis, no topology...
- large definitions; proofs with many similar cases;
- syntactic techniques (induction); few deep mathematical concepts.

The POPLmark challenge

In 2005, Aydemir *et al.* challenged the POPL community:

How close are we to a world where every paper on programming languages is accompanied by an electronic appendix with machine-checked proofs?

Today, over 20% of the papers at POPL come with such an appendix.

Proof assistants

An interactive proof assistant offers:

- A formal **specification language**,
in which definitions are written and theorems are stated.
- A set of **commands** for building proofs,
either automatically or interactively.
- Often, an independent, automated **proof checker**,
so the above commands do not have to be trusted.

Barendregt and Wiedijk,
The Challenge of Computer Mathematics, 2005.

Popular proof assistants include Rocq, Agda, HOL4, Isabelle/HOL...

Computations and functions

Rocq offers a pure functional programming language in the style of ML,
with recursive functions and pattern-matching.

```
Fixpoint factorial (n: nat) :=
  match n with
  | 0   => 1
  | S p => n * factorial p
  end.
```

```
Fixpoint concat (A: Type) (xs ys: list A) :=
  match xs with
  | nil => ys
  | x :: xs => x :: concat xs ys
  end.
```

The language is total: all functions terminate. This is enforced by
requiring every recursive call to be decreasing w.r.t. the subterm ordering.

Mathematical logic

Propositions can be expressed in this language. They have type **Prop**.

Definition divides (a b: N) := **exists** n: N, b = n * a.

Theorem factorial_divisors:

forall n i, 1 <= i <= n -> divides i (factorial n).

Definition prime (p: N) :=

p > 1 /\ (**forall** d, divides d p -> d = 1 \vee d = p).

Theorem Euclid:

forall n, **exists** p, p >= n /\ prime p.

The standard logical connectives and quantifiers are available.

Inductive types

An **inductive type** is a data type.

It is equipped with a finite number of **constructors**.

Its inhabitants are generated by repeated application of the constructors.

```
Inductive nat: Type :=
| 0: nat
| S: nat -> nat.
```

```
Inductive list: Type -> Type :=
| nil: forall A, list A
| cons: forall A, A -> list A -> list A.
```

E.g., the inhabitants of nat are 0, S 0, S (S 0), etc.

This is well suited to describing the **syntax** of a programming language.

Inductive predicates

An **inductive predicate** is equipped with a finite number of constructors, and is generated by repeated application of the constructors.

```
Inductive even: nat -> Prop :=
| even_zero:
  even 0
| even_plus_2:
  forall n, even n -> even (S (S n)).
```

On paper, this is typically written in the form of inference rules:

$$\frac{}{0 \text{ is even}} \qquad \frac{n \text{ is even}}{S(S n) \text{ is even}}$$

The inhabitants of the type `even n` can be thought of as **derivation trees** whose conclusion is `even n`.

Binding and α -equivalence

Most programming languages provide constructs that bind variables:

- function abstractions (in terms): $\lambda x.t$
- local definitions (in terms): *let* $x = t$ *in* t
- quantifiers (in types): $\forall \alpha. \alpha \rightarrow \alpha$

α -equivalence is a relation that allows renamings of bound variables, e.g.:

$$\lambda x. x + 1 \equiv_{\alpha} \lambda y. y + 1 \quad \forall \alpha. \alpha \text{ list} \equiv_{\alpha} \forall \beta. \beta \text{ list}$$

On paper, usually,

- the concepts of names and binding are not clearly defined;
- one confuses α -equivalence \equiv_{α} with equality $=$.

Representations of syntax

How should syntax with binding be mathematically defined,
on paper or in a proof assistant?

Several representations come to mind:

- de Bruijn notation – used in this course (de Bruijn, 1972);
- equivalence classes – the nominal approach (Pitts, 2006);
- (parametric) higher-order abstract syntax (Chlipala, 2008);
- the locally nameless representation (Charguéraud, 2009);
- the solutions of the POPLmark challenge (2005) involve 8 different representations, and there are more.

One should choose a representation for which the proof assistant has good support.

de Bruijn indices



A simple idea: don't use names.

Instead, use pointers from variables back to their binding site.

A second idea: use relative pointers, encoded as natural integers.

- 0 denotes the nearest enclosing λ ,
i.e., the most recently bound variable;
- 1 denotes the next enclosing λ , and so on.

$\lambda x.x$ is $\lambda 0$.

$\lambda f.\lambda x.f\ x$ is $\lambda\lambda(1\ 0)$.

Why is this a good idea?

de Bruijn syntax has several strengths:

- it is easily defined;
- it is **inductive** – terms are trees, no quotient is required;
- it is **canonical** – α -equivalence is just equality.

Its drawbacks are well-known, too:

- terms are more difficult to read;
- definitions and theorems can seem difficult to read and write
– mostly a matter of **habit**?

λ -terms in de Bruijn's notation

The syntax of λ -calculus is simple:

$$t ::= x \mid \lambda t \mid t t \quad \text{where } x \in \mathbb{N}$$

In Rocq:

```
Inductive term :=
| Var: nat -> term
| Lam: term -> term
| App: term -> term -> term.
```

Suggested exercises

Exercise: In OCaml, implement [conversions](#) between the nominal representation and de Bruijn's representation, both ways.

Exercise: In OCaml, implement an exhaustive [enumeration](#) of the λ -terms of size s and with at most n free variables. (Let variables have size 0; let λ -abstractions and applications contribute 1.)

Exercise: Use this exhaustive enumeration to [test](#) that the above conversions are inverses of each other.

Substitution



— *Substitution is the éminence grise of the λ -calculus.*
Abadi, Cardelli, Curien, Lévy, **Explicit substitutions**, 1990.

Substitutions

Let a substitution σ be a total function of variables \mathbb{N} to terms \mathbb{T} .

It can also be thought of as an infinite sequence $\sigma(0) \cdot \sigma(1) \cdot \dots$

Let id be the identity substitution: $id(x) = x$.

- $0 \cdot 1 \cdot 2 \cdot \dots$

Let $+i$ be the lift substitution: $(+i)(x) = x + i$.

- $i \cdot (i + 1) \cdot (i + 2) \cdot \dots$

Let $t \cdot \sigma$ be the cons substitution that maps 0 to t and $x + 1$ to $\sigma(x)$.

- $t \cdot \sigma(0) \cdot \sigma(1) \cdot \dots$

id can in fact be viewed as sugar for $0 \cdot (+1)$.

Substitution application and composition

Can we define $t[\sigma]$, the **application** of the substitution σ to the term t ?

It should satisfy the following laws:

$$\begin{aligned}x[\sigma] &= \sigma(x) \\ (\lambda t)[\sigma] &= ? \\ (t_1 \ t_2)[\sigma] &= t_1[\sigma] \ t_2[\sigma]\end{aligned}\quad \text{where } \uparrow\sigma \text{ stands for } 0 \cdot (\sigma ; +1)$$

$\uparrow\sigma$ can be understood as “the substitution σ , moved under one binder”.

The (left-to-right) **composition** of two substitutions $\sigma_1 ; \sigma_2$ should satisfy:

$$(\sigma_1 ; \sigma_2)(x) = (\sigma_1(x))[\sigma_2]$$

Unfortunately, these equations are mutually recursive,
so they are not accepted as a valid definition
by a proof assistant such as Rocq.

Work-around (1): define lifting first

Mechanization

Rocq in a
nutshellSyntax with
bindersde Bruijn
Nominal

To work around this problem, one first defines $t[\sigma]$ in the special case where σ is $\uparrow^i (+1)$.

This yields a function of two arguments, [lifting](#), written $\uparrow^i t$:

$$\begin{aligned}\uparrow^i(x) &= x + i \\ \uparrow^i(\lambda t) &= \lambda(\uparrow^{i+1} t) \\ \uparrow^i(t_1 t_2) &= \uparrow^i t_1 \uparrow^i t_2\end{aligned}$$

This definition is accepted by a proof assistant. (Why?)

$\uparrow^i t$ can be understood as “the term t , moved under i binders”.

Work-around (2): define substitution next

Then, one can define $\uparrow\sigma$ as the substitution that maps 0 to 0 and that maps $x + 1$ to $\uparrow^1(\sigma(x))$.

Exercise: Check that $\uparrow\sigma$ is equal to $0 \cdot (\sigma ; +1)$ as desired.

Finally, one defines substitution:

$$\begin{aligned}x[\sigma] &= \sigma(x) \\(\lambda t)[\sigma] &= \lambda(t[\uparrow\sigma]) \\(t_1 \ t_2)[\sigma] &= t_1[\sigma] \ t_2[\sigma]\end{aligned}$$

This definition is accepted by a proof assistant. (Why?)

de Bruijn algebra

The following equations are **valid**, that is, universally true:

$$\begin{array}{ll} (\lambda t)[\sigma] = \lambda(t[0 \cdot (\sigma; +1)]) & id; \sigma = \sigma \\ (t_1 \ t_2)[\sigma] = t_1[\sigma] \ t_2[\sigma] & \sigma; id = \sigma \\ 0[t \cdot \sigma] = t & (\sigma_1; \sigma_2); \sigma_3 = \sigma_1; (\sigma_2; \sigma_3) \\ (+1); (t \cdot \sigma) = \sigma & (t \cdot \sigma_1); \sigma_2 = t[\sigma_2] \cdot (\sigma_1; \sigma_2) \end{array}$$

Furthermore, these equations form a **complete theory**: that is, if an equation (which involves just the above concepts) is valid then it is a **logical consequence** of the above equations (**Schäfer et al., 2015**).

Schäfer et al. also prove that validity is **decidable**.

de Bruijn algebra

Decidability means that the machine can answer questions for us.

Does $t[id] = t$ hold? Yes.

Does $t[\sigma_1][\sigma_2] = t[\sigma_1 ; \sigma_2]$ hold? Yes.

Does $t[+1][u \cdot id] = t$ hold? Yes.

- In nominal style, this would be written: $t[u/x] = t$ if $x \notin fv(t)$.
- de Bruijn style is less familiar, but has no side condition.

And so on, and so forth.

For proofs of the first two equations above, see Schäfer et al., Fact 6.

We do not really care about these proofs – a machine can find them.

Rocq tactics for de Bruijn algebra

The Rocq library `AutoSubst` offers two tactics:

- `autosubst` proves an equation between terms or substitutions;
- `asimpl` simplifies a goal in which a term or substitution appears.

I use `AutoSubst` version 1 because it suffices for my needs.

λ -terms with AutoSubst

The syntax of λ -calculus can be declared as follows:

```
Inductive term :=  
| Var: var -> term  
| Lam: {bind term} -> term  
| App: term -> term -> term.
```

AutoSubst defines `var` as a synonym for `nat`
and `{bind term}` as a synonym for `term`.

AutoSubst defines substitution application, composition, etc., for us.

See [DemoSyntaxReduction](#).

AutoSubst key notations

$t \cdot \sigma$	<code>t .: sigma</code>	substitution “cons”
$+i$	<code>ren (+i)</code>	the substitution $+i$
id	<code>ids</code>	the identity substitution
$t[\sigma]$	<code>t.[sigma]</code>	substitution application
$\sigma_1 ; \sigma_2$	<code>sigmal >> sigma2</code>	substitution composition
$\uparrow \sigma$	<code>up sigma</code>	moving a substitution under a binder
$\uparrow^n \sigma$	<code>upn n sigma</code>	moving a substitution under n binders
$t.[u \cdot id]$	<code>t.[u/]</code>	substituting u for 0 in t

Reading “lift” as an end-of-scope mark

Suppose we are building a term in a context where n variables exist and we wish to refer to a subterm t that was built in a context where one of our variables, say x , where $0 \leq x < n$, did not exist.

We cannot just refer to t . Instead, we must use $t[\uparrow^x (+1)]$, that is, $\uparrow^x t$.

Is this ugly low-level index arithmetic?

Yes and no: one can read it as an end-of-scope mark for the variable x .

One can adopt a nicer meta-level notation for it, say “*unbind x in t*”.

The equation discussed earlier is then written $(\text{unbind } 0 \text{ in } t)[u \cdot id] = t$. That is, replacing 0 with u in a term where 0 is not in scope has no effect.

A related, object-level end-of-scope construct, “abdma”, has been studied by Hendriks and van Oostrom (2003).

Calculi of explicit substitutions

Similarly, we have viewed substitution application as a meta-level operation. There is no syntax for it in the λ -calculus.

In the $\lambda\sigma$ -calculus, however, there is [syntax for substitutions](#) and [substitution application](#), and a set of [small-step reduction rules](#) that explain how substitutions interact with λ -abstractions and applications.

Abadi, Cardelli, Curien, Lévy, [Explicit substitutions](#), 1990.

Curien, Hardin, Lévy, [Confluence properties of weak and strong calculi of explicit substitutions](#), 1992.

Binding and α -equivalence

Most programming languages provide constructs that bind variables:

- function abstractions (in terms): $\lambda x.t$
- local definitions (in terms): $\text{let } x = t \text{ in } t$
- quantifiers (in types): $\forall \alpha. \alpha \rightarrow \alpha$

α -equivalence is a relation that allows renamings of bound variables, e.g.:

$$\lambda x. x + 1 \equiv_{\alpha} \lambda y. y + 1 \quad \forall \alpha. \alpha \text{ list} \equiv_{\alpha} \forall \beta. \beta \text{ list}$$

α -equivalence, a relation on trees, can be defined as follows (Pitts, 2006):

$$\frac{y \notin \text{fv}(\lambda x.t) \quad u = \binom{x}{y} t}{\lambda x.t \equiv_{\alpha} \lambda y.u}$$

where $\binom{x}{y}$ swaps all occurrences of the names x and y in the term t .

Implicit α -equivalence

On paper, it is customary to confuse α -equivalence \equiv_α with equality $=$.

This plays a role, for instance, in the definition of System F ,
where types contain universal quantifiers.

This is the traditional rule for type-checking a function application:

$$\frac{\Gamma \vdash e_1 : \tau \rightarrow \tau' \quad \Gamma \vdash e_2 : \tau}{\Gamma \vdash e_1 e_2 : \tau'}$$

If α -equivalence was explicit, the rule would be written as follows:

$$\frac{\Gamma \vdash e_1 : \tau \rightarrow \tau' \quad \Gamma \vdash e_2 : \tau_2 \quad \tau \equiv_\alpha \tau_2}{\Gamma \vdash e_1 e_2 : \tau'}$$

In simply-typed λ -calculus, this issue does not arise, as there are no
quantifiers in types: α -equivalence and equality of types coincide.

Explicit α -equivalence

In principle, one should distinguish between:

- trees versus equivalence classes of trees;
- equality = versus α -equivalence \equiv_α .

A term is an equivalence class of trees.

This sounds easy enough, but leads to subtleties when defining mathematical functions that consume or produce terms... such as:

- program transformations, which produce and consume syntax trees;
- proofs, which produce and consume derivation trees.

Functions on equivalence classes

To define a **function** f from \mathbb{T}/\equiv_α to \mathbb{T}/\equiv_α ,
it suffices to first define a **relation** F between \mathbb{T} and \mathbb{T} ,
and to require two conditions:

- every tree is α -equivalent to some tree in the domain of F :

$$\forall t \in \mathbb{T} \quad \exists t', u' \in \mathbb{T} \quad t \equiv_\alpha t' \wedge t' F u'$$

- note: the domain of F need not be \mathbb{T}
- “without loss of generality, let us assume that x does not occur in ...”

- F is **compatible** with α -equivalence:

$$\forall t, t', u, u' \in \mathbb{T} \quad t F u \wedge t' F u' \wedge t \equiv_\alpha t' \Rightarrow u \equiv_\alpha u'$$

- note: F need not be deterministic (single-valued)
- nondeterminism is fine as long as all choices yield α -eq. results
- “let us pick a name x outside of ...”

Free variables

This is the classic definition of the set of the **free variables** of a tree:

$$\begin{aligned}fv(x) &= \{x\} \\fv(\lambda x.t) &= fv(t) \setminus \{x\} \quad - \text{no requirement on } x \\fv(t_1 t_2) &= fv(t_1) \cup fv(t_2)\end{aligned}$$

It is a total function from \mathbb{T} to sets of names.

Condition 1 is vacuously satisfied (the relation is defined everywhere).

Condition 2 requires checking the following equality:

$$fv(\lambda x.t) = fv(\lambda y.(^x_y)t) \quad \text{where } y \notin fv(\lambda x.t)$$

This follows from the fact that fv is **equivariant**,
that is, fv commutes with every permutation π :

$$fv(\pi t) = \pi fv(t)$$

and from the fact that neither x nor y appear in the set $fv(\lambda x.t)$.

Thus, fv gives rise to a total function from \mathbb{T}/\equiv_α to sets of names.

Capture-avoiding substitution

This is the classic definition of capture-avoiding substitution:

$$\begin{aligned}x[u/x] &= u \\y[u/x] &= y && \text{if } y \neq x \\(\lambda z. t)[u/x] &= \lambda z. t[u/x] && \text{if } z \notin fv(u) \cup \{x\} && - \text{ avoid capture!} \\(t_1 \ t_2)[u/x] &= t_1[u/x] \ t_2[u/x]\end{aligned}$$

It is a partial function from \mathbb{T} to \mathbb{T} .

Condition 1 holds, as only a finite number of choices for z are forbidden.

Condition 2 requires checking:

$$\lambda z. t[u/x] \equiv_{\alpha} \lambda z'. t'[u/x] \quad \text{where } z, z' \notin fv(u) \cup \{x\} \text{ and } \lambda z. t \equiv_{\alpha} \lambda z'. t'$$

which follows, again, from the fact that substitution is equivariant.

Thus, this gives rise to a total function from $\mathbb{T}/\equiv_{\alpha}$ to $\mathbb{T}/\equiv_{\alpha}$.

Naïve substitution

Naïve substitution does not have the side condition $z \notin \text{fv}(u) \cup \{x\}$.

It is a total function from \mathbb{T} to \mathbb{T} ,
but does not satisfy condition 2,
hence does not give rise to a function from \mathbb{T}/\equiv_α to \mathbb{T}/\equiv_α .

$$(\lambda y. x + y)[2 \times y/x] = \lambda y. 2 \times y + y \quad - \text{naïve}$$

$$(\lambda y. x + y)[2 \times y/x] =$$

$$(\lambda z. x + z)[2 \times y/x] = \lambda z. 2 \times y + z \quad - \text{capture-avoiding}$$

The nominal approach in proof assistants

The nominal approach is prevalent in informal (paper) proofs.

It is implemented in Nominal Isabelle (Urban, 2008).

- Urban and Narboux (2008) present typical proofs about operational semantics.

It is not well supported in Rocq, perhaps for engineering reasons.

- Cohen (2013) shows how to use quotients in Rocq (when they exist) and how to construct them (up to certain axioms or hypotheses).