Making the stack explicit: the continuation-passing style transformation

MPRI 2.4

François Pottier

2024

Motivation

What if a program transformation could:

- ensure that every function call is a tail call and the stack is explicit, so the code is no longer really recursive, but iterative;
- make the evaluation order explicit in the code, so that it does not depend on the ambient strategy (CBN / CBV);
- eliminate the apparent redundancy between calls and returns, by exploiting solely function calls – functions never return!
- suggest extending the λ-calculus with control operators?

The continuation-passing style transformation does all this.

Motivation

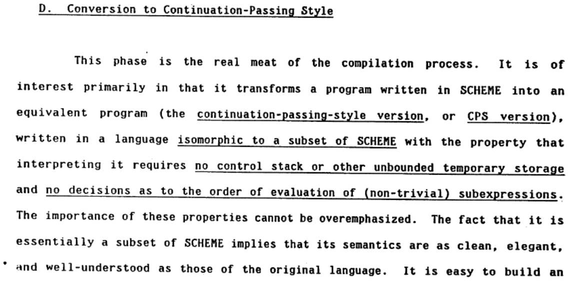

D. Conversion to Continuation-Passing Style

This phase is the real meat of the compilation process. It is of interest primarily in that it transforms a program written in SCHEME into an equivalent program (the continuation-passing-style version, or CPS version), written in a language isomorphic to a subset of SCHEME with the property that interpreting it requires no control stack or other unbounded temporary storage and no decisions as to the order of evaluation of (non-trivial) subexpressions. The importance of these properties cannot be overemphasized. The fact that it is essentially a subset of SCHEME implies that its semantics are as clean, elegant, and well-understood as those of the original language. It is easy to build an

Steele, RABBIT: a compiler for SCHEME, 1978.

A direct-style interpreter

Recall our environment-based interpreter for call-by-value λ-calculus:

```
let rec eval (e : cenv) (t : term) : cvalue =
  match t with
  | Var x ->
      lookup e x
  | Lam t ->
      Clo (t, e)
  | App (t1, t2) ->
      let cv1 = eval e t1 in
      let cv2 = eval e t2 in
      let Clo (u1, e') = cv1 in
      eval (cv2 :: e') u1
```

This is an OCaml transcription, without a fuel parameter.

A continuation-passing style interpreter

Instead of returning a value,

```
let rec eval (e : cenv) (t : term) : cvalue =
  ...
```

let's pass this value to a continuation that we get as an argument:

```
let rec evalk (e : cenv) (t : term) (k : cvalue -> 'a) : 'a =
  ...
```

Exercise (in class): write `evalk`. (See `EvalCBVExercise`.)

A continuation-passing style interpreter

```
let rec evalk (e : cenv) (t : term) (k : cvalue -> 'a) : 'a =
  match t with
  | Var x ->
      k (lookup e x)
  | Lam t ->
      k (Clo (t, e))
  | App (t1, t2) ->
      evalk e t1 (fun cv1 ->
      evalk e t2 (fun cv2 ->
      let Clo (u1, e') = cv1 in
      evalk (cv2 :: e') u1 k))
```

Instead of returning a value, pass it to `k`.

Instead of sequencing computations via `let`, nest continuations.

A continuation-passing style interpreter

To run the interpreter, start it with the identity continuation:

```
let eval (e : cenv) (t : term) : cvalue =
  evalk e t (fun cv -> cv)
```

Correctness of the CPS interpreter

The continuation-passing style interpreter is “obviously” correct.

Exercise: define `evalk` in Coq (with fuel) and prove it equivalent to the direct-style interpreter: `evalk n e t k = k (eval n e t)`.

Properties of the interpreter

What is special about this interpreter?

- Every call to `evalk` is a tail call.
- Every call to a continuation `k` is a tail call.

Tail calls

A call $g\ x$ is a tail call if it is the "last thing" that the calling function does...

More formally,

$v ::= x \mid \lambda x.tt$	values
$tt ::=$	terms in tail position
$\mid\ v$	
$\mid\ nt\ nt$	– a tail call
$\mid\ let\ nt\ in\ tt$	
$\mid\ if\ nt\ then\ tt\ else\ tt$	
$nt ::=$	terms not in tail position
$\mid\ v$	
$\mid\ nt\ nt$	– not a tail call
$\mid\ let\ nt\ in\ nt$	
$\mid\ if\ nt\ then\ nt\ else\ nt$	

This can be understood as the description of a top-down computation that assigns a Boolean flag ("tail" or "non-tail") to every subterm.

Verified tail calls

OCaml allows us to verify that these are indeed tail calls:

```
let rec evalk (e : cenv) (t : term) (k : cvalue -> 'a) : 'a =
  match t with
  | Var x ->
      (k[@tailcall]) (lookup e x)
  | Lam t ->
      (k[@tailcall]) (Clo (t, e))
  | App (t1, t2) ->
      (evalk[@tailcall]) e t1 (fun cv1 ->
      (evalk[@tailcall]) e t2 (fun cv2 ->
      let Clo (u1, e') = cv1 in
      (evalk[@tailcall]) (cv2 :: e') u1 k))
```

A nice feature (though with somewhat ugly syntax).

Properties of the interpreter

Tail calls are compiled by OCaml to jumps.

Thus, tail-recursive functions are compiled by OCaml to loops.

Steele, Lambda: the ultimate GOTO, 1977.

Thus, the CPS interpreter is not truly recursive: it is iterative.

It uses constant space on OCaml's implicit stack.

Wait! Does the interpreter really not need a stack any more?

- Of course it does need a stack.
- The continuation, allocated in the OCaml heap, serves as a stack.

A defunctionalized CPS interpreter

To better see the structure of the continuation,
let us defunctionalize the CPS interpreter.

Reynolds, Definitional interpreters
for programming languages, 1972 (1998).

Reynolds, Definitional interpreters revisited, 1998.

Defunctionalization (reminder)

Steps:

- Identify the sites where closures are allocated,
  that is, where anonymous functions are built.
- Compute, at each site, the free variables of the anonymous function.
- Introduce an algebraic data type of closures.
- Transform the code:
  - replace anonymous functions with constructor applications,
  - replace function applications with calls to `apply`,
  - and define `apply`.

Exercise (in class): defunctionalize the CPS interpreter. (`EvalCBVExercise`.)

A defunctionalized CPS interpreter

There are three sites where an anonymous continuation is built.

We name them and compute their free variables.

This leads to the following algebraic data type of continuations:

```
type kont =
  | AppL of { e: cenv; t2: term; k: kont }
  | AppR of {       cv1: cvalue; k: kont }
  | Init
```

What data structure is this? A linked list. A heap-allocated stack.

In fact, it is a (call-by-value) evaluation context:

$$E ::= E[[]\ \ t_2[e]] \mid E[v_1\ \ []] \mid []$$

It is a zipper, a path from the context's hole up to the root of a term.

Huet, The Zipper, 1997.

A defunctionalized CPS interpreter

We transform the interpreter's main function:

```
let rec evalkd (e : cenv) (t : term) (k : kont) : cvalue =
  match t with
  | Var x ->
      apply k (lookup e x)
  | Lam t ->
      apply k (Clo (t, e))
  | App (t1, t2) ->
      evalkd e t1 (AppL { e; t2; k })
```

To evaluate $t_1\ t_2$, the interpreter pushes information on the stack, then jumps straight to evaluating t_1.

A defunctionalized CPS interpreter

`apply` interprets continuations as functions of values to values:

```
and apply (k : kont) (cv : cvalue) : cvalue =
  match k with
  | AppL { e; t2; k } ->
      let cv1 = cv in
      evalkd e t2 (AppR { cv1; k })
  | AppR { cv1; k } ->
      let cv2 = cv in
      let Clo (u1, e') = cv1 in
      evalkd (cv2 :: e') u1 k
  | Init ->
      cv
```

It pops the top stack frame and decides what to do, based on it.

A defunctionalized CPS interpreter

To run the interpreter, start it with the identity continuation:

```
let eval e t =
  evalkd e t Init
```

An abstract machine

We have reached an abstract machine, a simple iterative interpreter which maintains a few data structures:

- a code pointer: the term `t`,
- an environment `e`,
- a stack, or continuation `k`.

In fact, we have mechanically rediscovered the CEK machine.

Felleisen and Friedman,
Control operators, the SECD machine, and the λ-calculus, 1987.

Sig Ager, Biernacki, Danvy and Midtgaard,
A Functional Correspondence between Evaluators and Abstract Machines, 2003.

Re-discovering other abstract machines

Exercise: start with a call-by-name interpreter and follow an analogous process to rediscover Krivine's machine.

The solution is in `EvalCBNCPS`.

There once was a man named Krivine
Who invented a wond'rous machine.
It pushed and it popped
On abstractions it stopped;
That lean mean machine from Krivine.
— *Mitchell Wand*

Krivine, A call-by-name lambda-calculus machine, (1985) 2007.

Formulations of the CPS transformation

There are many variants of the CPS transformation,
and sometimes many formulations of a single variant.

Let us look at the simplest formulation: Fischer and Plotkin's.

Fischer, Lambda-Calculus Schemata, (1972) 1993.

Plotkin, Call-by-name, call-by-value and the λ-calculus, 1975.

Definition of the CBV CPS transformation

A term is translated to a function of a continuation k to an answer.

$$\begin{aligned}
[\![x]\!] &= \lambda k.\ k\ x \\
[\![\lambda x.t]\!] &= \lambda k.\ k\ (\lambda x.[\![t]\!]) \\
[\![t_1\ t_2]\!] &= \lambda k.\ [\![t_1]\!]\ (\lambda x_1.\ [\![t_2]\!]\ (\lambda x_2.\ x_1\ x_2\ k)) \\
[\![let\ x = t_1\ in\ t_2]\!] &= \lambda k.\ [\![t_1]\!]\ (\lambda x.\ [\![t_2]\!]\ k)
\end{aligned}$$

A value $\lambda x.t$ is translated to a function of two arguments $\lambda x.\lambda k.\ldots$.

Definition of the CBV CPS transformation

One avoids some redundancy by defining two mutually recursive functions, namely the translation of values $(\!|v|\!)$:

$$
\begin{aligned}
(\!|x|\!) &= x \\
(\!|\lambda x.t|\!) &= \lambda x.[\![t]\!]
\end{aligned}
$$

and the translation of terms $[\![t]\!]$:

$$
\begin{aligned}
[\![v]\!] &= \lambda k.\ k\ (\!|v|\!) \\
[\![t_1\ t_2]\!] &= \lambda k.\ [\![t_1]\!]\ (\lambda x_1.\ [\![t_2]\!]\ (\lambda x_2.\ x_1\ x_2\ k)) \\
[\![\mathit{let}\ x = t_1\ \mathit{in}\ t_2]\!] &= \lambda k.\ [\![t_1]\!]\ (\lambda x.\ [\![t_2]\!]\ k)
\end{aligned}
$$

Indifference

In a transformed term, the right-hand side of every application is a value.

Therefore, its execution is indifferent to the choice
of a call-by-name or call-by-value evaluation strategy.

In other words, evaluation order is fully explicit in a transformed term.

The transformation on the previous slide fixes a call-by-value strategy:
it is the CBV CPS transformation.

It can serve as an encoding of call-by-value into call-by-name,
thus answering a question raised in week 1.

Exercise (recommended): Define the CBN CPS transformation.

Stacklessness

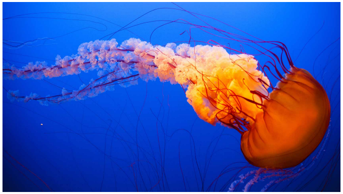

In a transformed term, every call is a tail call.

Therefore, reduction under a context is not required.

That is, execution does not require a stack.

We could (but won't) give a (small-step, substitution-based) semantics that takes indifference and stacklessness into account.

Exercise: Propose such a semantics. Prove that, when executing a CPS-transformed term, it is equivalent to the standard semantics.

Effect of the transformation of types

How are types transformed?

A value of type T is translated to a value of type $(\!|T|\!)$.

A computation of type T is translated to a value of type $[\![T]\!]$.

$$
\begin{aligned}
(\!|\alpha|\!) &= \alpha \\
(\!|T_1 \rightarrow T_2|\!) &= (\!|T_1|\!) \rightarrow [\![T_2]\!] \\
[\![T]\!] &= ((\!|T|\!) \rightarrow A) \rightarrow A
\end{aligned}
$$

The type A, known as the answer type, is arbitrary and fixed.

One may take A to be the empty type 0. Then, $[\![T]\!]$ is $\neg\neg(\!|T|\!)$. The CPS transformation is known in logic as the double-negation translation.

Exercise (recommended): state and prove Type Preservation.

Effect of the transformation of types – refined

Could the transformation of types be made more precise in some sense?

$$[\![T]\!] = ((\!|T|\!) \rightarrow A) \rightarrow A$$

Every transformed term is in fact answer-type polymorphic:

$$[\![T]\!] = \forall A.((\!|T|\!) \rightarrow A) \rightarrow A$$

Furthermore, every transformed term invokes its continuation once:

$$[\![T]\!] = \forall A.((\!|T|\!) \rightarrow A) \multimap A$$

However, these properties are violated in the presence of control effects.

Thielecke, From control effects to typed continuation passing, 2003.

Semantic preservation

Plotkin (1975) proved semantic preservation,
based on a small-step simulation diagram.

This proof is complicated by the presence of administrative reductions.

A simpler approach is to use big-step semantics in the hypothesis:

Lemma (Semantic Preservation)

If $t \downarrow_{cbv} v$ and if w is a value, then $[\![t]\!]\; w \longrightarrow^{\star}_{cbv} w\; (\!|v|\!)$.

One should prove, in addition, that divergence is preserved.

Exercise (recommended): Prove this lemma.

Monadic intermediate form

If one just aims to make evaluation order explicit, CPS is overkill.

This transformation, too, achieves indifference:

$$\begin{aligned}
[\![x]\!] &= x \\
[\![\lambda x.t]\!] &= \lambda x.[\![t]\!] \\
[\![t_1\ t_2]\!] &= \mathit{let}\ x_1 = [\![t_1]\!]\ \mathit{in} \\
&\ \mathit{let}\ x_2 = [\![t_2]\!]\ \mathit{in} \\
&\ x_1\ x_2 \\
[\![\mathit{let}\ x = t_1\ \mathit{in}\ t_2]\!] &= \mathit{let}\ x = [\![t_1]\!]\ \mathit{in}\ [\![t_2]\!]
\end{aligned}$$

In a transformed term, the components of every application are values.

By further hoisting "*let*" out of the left-hand side of "*let*",
one gets administrative normal form.

Flanagan, Sabry, Felleisen, The essence
of compiling with continuations, 1993 (2003).

The CPS monad

The CPS transformation is a special case of the monadic transformation.

See Dagand’s lectures!

Some history

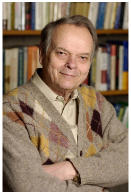

Continuations, and the CPS transformation, were independently discovered by many researchers during the 1960s.

John C. Reynolds, The discoveries of continuations, 1993.

Some history

The CPS transformation has been used in compilers.

Rabbit (Steele). SML/NJ.

Appel, Compiling with Continuations, 1992.

Today, heap-allocating the stack is considered too costly:

- bad locality;
- increased GC load;
- confuses the processor's built-in prediction of return addresses.

Yet, selective CPS transformations are used to compile effect handlers,

and some compilers use CPS as an intermediate form before coming back to direct style.

Kennedy, Compiling with continuations, continued, 2007.

Some history

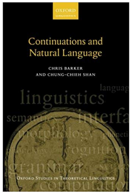
OXFORD LINGUISTICS
Continuations and Natural Language
CHRIS BARKER AND CHUNG-CHIEH SHAN
Oxford Studies in Theoretical Linguistics

Can λ-calculus and continuations explain the structure of speech?

Chris Barker,
Continuations and the nature of quantification, 2002.

Chris Barker and Chung-Chieh Shan,
Continuations and Natural Language, 2014.

A few things to remember

Continuations rule!

- The CPS transformation achieves several remarkable effects:
  - making the stack explicit;
  - making evaluation order explicit;
  - suggesting/explaining control operators.
- It plays a fundamental role in prog. language theory and in logic.
- Continuation-passing is also a useful programming technique.