

Introduction to Rust

Jacques-Henri Jourdan

January 25th, 2023

Low-level programming

Introduction to
Rust

Jacques-Henri
Jourdan

Some software requires a high level of control:

- over memory layout;
- over when memory is allocated and freed;
- over what computations are done and when...

The standard “answer” to these issues is C/C++.

These languages have shortcomings:

- very complicated semantics (pointer optimizations...),
- no safety guarantee,
- poor support for abstraction.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects



A language initially developed by Mozilla for rewriting parts of Firefox

High performance, high level of control with safe abstraction mechanisms



A language initially developed by Mozilla for rewriting parts of Firefox

High performance, high level of control with safe abstraction mechanisms

Zero-cost abstraction

Powerful abstraction mechanisms with no impact on performance:

- Sound type system
- Pointers: ownership + borrows with lifetimes
- Polymorphism with traits (\simeq type classes)

When the type system is not expressive enough, we can use **unsafe features** behind safe abstractions.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

[Basic types](#)[Aliasing control](#)[Ownership in Rust](#)[Borrows and lifetimes](#)[Generics & traits](#)[Iterators](#)[Closures](#)[Trait objects](#)

Why study Rust here?

Many concepts from type systems, adapted them to a new context:

- Polymorphism, type traits, closures, algebraic types...

New type system features:

- Ownership types, borrows, lifetimes,
- Concurrency: fine-grained tracking of thread-(un)safe types.

New challenges in type system meta-theory:

- ownership,
- unsafe code behind safe abstractions.

[Basic types](#)[Aliasing control](#)[Ownership in Rust](#)[Borrows and lifetimes](#)[Generics & traits](#)[Iterators](#)[Closures](#)[Trait objects](#)

What you will learn here

My goal is that in the following 5 weeks, you will:

- know how to write simple Rust programs;
- roughly understand how the type system works,
 - both formally and how it can be implemented;
- have a glimpse of how to formally study Rust:
 - How to use the type system to verify Rust programs?
 - How to build a logical relation to prove soundness?

Today: learn a bit of the language

Resources

Introduction to
Rust

Jacques-Henri
Jourdan

To learn Rust as a developer:

- Plenty of references here: <https://www.rust-lang.org/learn>
- Some of my favorites:
 - [Rust 101](#) by Ralf Jung: accessible tutorial for the core language;
 - Many examples of this course taken from this tutorial
 - [Rust by example](#): learn by example;
 - [The Rust Book](#): comprehensive, but long, “official” tutorial for Rust and some of its ecosystem.

Metatheory of the type system:

- RustBelt: Securing the Foundations of the Rust Programming Language. Ralf Jung, Jacques-Henri Jourdan, Robbert Krebbers, Derek Dreyer. POPL 2018.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

1 Basic types

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

2 Aliasing control

- Ownership in Rust
- Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

3 Generics & traits

- Iterators
- Closures
- Trait objects

The minimum of a sequence of integers

```
enum NumberOrNothing {
    Number(i32), Nothing
}

fn vec_min(vec: Vec<i32>) -> NumberOrNothing {
    let mut res = NumberOrNothing::Nothing;
    for el in vec {
        match res {
            NumberOrNothing::Nothing => {
                res = NumberOrNothing::Number(el);
            },
            NumberOrNothing::Number(n) => {
                let m = if n < el { n } else { el };
                res = NumberOrNothing::Number(m);
            }
        }
    }
    return res
}
```

A function for computing the maximum of a sequence of 32-bits integers.

- Parameters and return types need to be annotated.
- Types for local variables are inferred (most often).

The minimum of a sequence of integers

```
enum NumberOrNothing {  
    Number(i32), Nothing  
}
```

```
fn vec_min(vec: Vec<i32>) -> NumberOrNothing {  
    let mut res = NumberOrNothing::Nothing;  
    for el in vec {  
        match res {  
            NumberOrNothing::Nothing => {  
                res = NumberOrNothing::Number(el);  
            },  
            NumberOrNothing::Number(n) => {  
                let m = if n < el { n } else { el };  
                res = NumberOrNothing::Number(m);  
            }  
        }  
    }  
  
    return res  
}
```

Empty sequence
⇒ no defined minimum

We use an algebraic data type

- 32 bits integer, or nothing
- pattern matching
- constructors

The minimum of a sequence of integers

```
enum NumberOrNothing {
    Number(i32), Nothing
}

use NumberOrNothing::*;

fn vec_min(vec: Vec<i32>) -> NumberOrNothing {
    let mut res = Nothing;
    for el in vec {
        match res {
            Nothing => {
                res = Number(el);
            },
            Number(n) => {
                let m = if n < el { n } else { el };
                res = Number(m);
            }
        }
    }
    return res
}
```

Empty sequence
⇒ no defined minimum

We use an algebraic data type

- 32 bits integer, or nothing
- pattern matching
- constructors
 - In a specific namespace, can be imported

The minimum of a sequence of integers

```
enum NumberOrNothing {  
    Number(i32), Nothing  
}  
  
use NumberOrNothing;  
  
fn vec_min(vec: Vec<NumberOrNothing>) -> i32 {  
    let mut res = NumberOrNothing::Nothing;  
    for el in vec {  
        match res {  
            NumberOrNothing::Nothing =>  
                res = el;  
            NumberOrNothing::Number(n) =>  
                if n < el {  
                    res = el;  
                } else {  
                    res = Number(n);  
                }  
        }  
    }  
  
    res.  
}
```

A value of type `NumberOrNothing` is a **sequence of bytes** (tag followed by `i32` payload).

- Control on memory representation, like in C/C++.
- Drawback: values have **sizes**, which the compiler need to know.
- In higher-level like OCaml, Java, Python, Haskell, ... this is different:
 - complex values are boxed;
 - values always use **one** memory word;
 - unavoidable implicit indirection.

Using vec_min

```
impl NumberOrNothing {
    fn print(self) {
        match self {
            Nothing =>
                println!("The number is: <nothing>"),
            Number(n) =>
                println!("The number is: {}", n),
        };
    }
}

fn main() {
    let v = vec![18,5,7,2,9,27];
    let min = vec_min(v);
    min.print()
}
```

We can define **associated functions** for types:

- like sealed methods in OO languages,
- no more than a function, but without polluting the global namespace.

Records, tuples, arrays, newtype

We (of course!) have other convenient ways to build types:

- Tuples (`T1, T2, ...`):
 - construction: (`e1, e2, ...`)
 - projections: `t.0, t.1, ...`
- Records declared by `struct R { f1: T1, f2: T2, ... }`
 - construction: `R { f1: e1, f2: e2, ... }`
 - projections: `r.f1, r.f2, ...`
- Newtype (as in Haskell) declared by `struct NT(T)`
 - construction: `NT(e)`
 - projection: `nt.0`
- Fixed-length arrays [`T; 42`]
 - construction: `[e; 42]`
 - access: `a[i]`

Records, tuples, arrays, newtype

We (of course!) have other convenient ways to build types:

- Tuples (`T1, T2, ...`):
 - construction: (`e1, e2, ...`)
 - projections: `t.0, t.1, ...`
- Records declared by `struct R { f1: T1, f2: T2, ... }`
 - construction: `R { f1: e1, f2: e2, ... }`
 - projections: `r.f1, r.f2, ...`
- Newtype (as in Haskell) declared by `struct NT(T)`
 - construction: `NT(e)`
 - projection: `nt.0`
- Fixed-length arrays [`T; 42`]
 - construction: `[e; 42]`
 - access: `a[i]`

Again, records, tuples, enum and arrays are stored with **no implicit pointer indirection**

- Values of tuples, records, arrays and enums types are the concatenation of their components.

1 Basic types

Basic types

2 Aliasing control

- Ownership in Rust
- Borrows and lifetimes

Aliasing control

Ownership in Rust

Borrows and lifetimes

3 Generics & traits

- Iterators
- Closures
- Trait objects

Generics & traits

Iterators

Closures

Trait objects

Ownership: counter example in C++

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
void foo(std::vector<int> v) {  
    int *first = &v[0];  
    v.push_back(42);  
    *first = 1337;  
}
```

Where is the bug?

Ownership: counter example in C++

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
void foo(std::vector<int> v) {  
    int *first = &v[0];  
    v.push_back(42);  
    *first = 1337;  
}
```

- `push_back` may **reallocate** the vector,
 - invalidating inner pointers such as `first`.
- Somewhat perverse bug:
 - most of the time, `push_back` does not reallocate the vector.
- We mutated memory through two aliased pointers, `first` and `v`.

Bugs caused by aliasing

Iterator invalidation:

- Mutating a data structure invalidates its iterators.
 - In C++, the rule is complicated (depends on where you iterate and what mutation...).
 - Common source of bugs in C++/Java/...
 - Often undetected by the library: unexpected consequences.
 - Core of the problem: **mutation and aliasing at the same time**.

Double free:

- Can lead to state inconsistency of the memory management library.
- Another form of simultaneous **mutation (free) and aliasing**.

Data races:

- Particularly difficult to test/debug.
- By definition: **mutation and aliasing**.

In Rust, we maintain the invariant: mutation XOR aliasing.

If we want to mutate, we need to have the **unique alias**.

- We are the **owner** of the memory location.
- We view memory as a **resource** which cannot be duplicated.

This also applies to **freeing** memory.

- There is no GC in Rust (better control of resources, better performances).

Ownership in vec_min

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
impl NumberOrNothing {
    fn print(self) { ... }

fn vec_min(vec: Vec<i32>) -> NumberOrNothing {
    ...
}

fn main() {
    let mut v = vec![18,5,7,2,9,27];
    vec_min(v).print();

}
```

Ownership in vec_min

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
impl NumberOrNothing {
    fn print(self) { ... }

fn vec_min(vec: Vec<i32>) -> NumberOrNothing {
    ...
}

fn main() {
    let mut v = vec![18,5,7,2,9,27];
    vec_min(v).print();
    v.push(42);
    vec_min(v).print();
}
```

Ownership in vec_min

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
impl NumberOrNothing {
    fn print(self) { ... }

}

fn vec_min(vec: Vec<i32>) -> NumberOrNothing {
    ...
}

fn main() {
    let mut v = vec![18, 5, 7, 2, 9, 27];
    vec_min(v).print();
    v.push(42);
    vec_min(v).print();
}

error[E0382]: borrow of moved value: 'v'
--> src/main.rs:37:3
   |
35 |     let mut v = vec![18,5,7,2,9,27];
   |     ...
36 |     vec_min(v).print();
   |             - value moved here
37 |     v.push(42);
   |     ^^^^^^^^^ value borrowed here after move
```

```
impl NumberOrNothing {
    fn print(self) { ... }
}

fn vec_min(vec: Vec<i32>) -> NumberOrNothing {
    ...
}

fn main() {
    let mut v = vec![18,5,7,2,9,27];
    vec_min(v).print();
    v.push(42);
    vec_min(v).print();
}
```

Passing `v` as a parameter to `vec_min` triggered a move.

- We lost its ownership. We can no longer use it.
- That's fortunate, because `vec_min` automatically freed the vector.
 - RAII: destructor is called as soon as variable leaves its scope (except if already moved).

The iconic ownership type: Box<T>

Box<T> is the type of owned pointers to T.

- Creation function:

```
fn new(x: T) -> Box<T>
```

- Move of x to dynamically allocated memory (i.e., malloc).
- Transfer of ownership.
- Can be dereferenced (for read/write): *p.
- RAII: Automatically freed when leaves the scope.

Example of use: singly linked list:

```
type List = Option<Box<Node>>;
struct Node { elem: i32, next: List, }
```

Exception to ownership tracking

Is ownership tracking a good thing for all variables?

Exception to ownership tracking

No! Some values are **bare data**, and can be copied freely.

- Examples: `i32`, `(i64, u64)`, `bool`, ...

We say that these types **implement the `Copy` trait**

- Tells the compiler to disable ownership tracking for values of these types.

By default, types declared with `struct` or `enum` are not `Copy`.

- Because they could be used to represent resources not known by the compiler.
- To make it so, this is easy:

```
#[derive(Copy, Clone)]
enum NumberOrNothing {
    Number(i32), Nothing
}
```

- (More explanations on traits later.)

Shared borrows

Variant of `vec_min` to avoid the move:

```
fn vec_min(vec: &Vec<i32>) -> NumberOrNothing {
    let mut res = Nothing;
    for el in vec {
        match res {
            Nothing => { res = Number(*el); },
            Number(n) => {
                let m = if n < *el { n } else { *el };
                res = Number(m);
            }
        }
    }
    return res
}

fn main() {
    let mut v = vec![18,5,7,2,9,27];
    vec_min(&v).print();
    v.push(42);
    vec_min(&v).print();
}
```

Instead of passing the ownership of the vector we pass a **shared borrow** of it.

- At runtime: a pointer to `v`
- Here, the shared borrow is only valid for the duration of the function.
- `&T` implements `Copy`.
- `&T` does not allow mutation (usually... see next course).
- aka. **immutable borrow**, **immutable reference**.

Shared borrows

Variant of `vec_min` to avoid the move:

```
fn vec_min(vec: &Vec<i32>) -> NumberOrNothing {
    let mut res = Nothing;
    for el in vec {
        match res {
            Nothing => { res = Number(*el); },
            Number(n) => {
                let m = if n < *el { n } else { *el };
                res = Number(m);
            }
        }
    }
    return res
}

fn main() {
    let mut v = vec![18,5,7,2,9,27];
    vec_min(&v).print();
    v.push(42);
    vec_min(&v).print();
}
```

What is the type of `el`? Why?

Shared borrows

Variant of `vec_min` to avoid the move:

```
fn vec_min(vec: &Vec<i32>) -> NumberOrNothing {
    let mut res = Nothing;
    for el in vec {
        match res {
            Nothing => { res = Number(*el); },
            Number(n) => {
                let m = if n < *el { n } else { *el };
                res = Number(m);
            }
        }
    }
    return res
}

fn main() {
    let mut v = vec![18,5,7,2,9,27];
    vec_min(&v).print();
    v.push(42);
    vec_min(&v).print();
}
```

When we access elements of a shared borrow of a vector (e.g., by iterating with a `for` loop), we get **shared borrows of the content**, because, in general, we cannot transfer ownership out of the vector.

Thus we need to **dereference `el!`**

Unique borrows

Introduction to
Rust

Jacques-Henri
Jourdan

```
fn vec_inc(vec: &mut Vec<i32>) {  
    for el in vec {  
        *el += 1  
    }  
  
    fn main() {  
        let mut v = vec![18,5,7,2,9,27];  
        vec_inc(&mut v);  
        v.push(42); /* ... */  
    }  
}
```

Unique borrows are similar to shared borrows.

Differences:

- Allows mutation.
- ...

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Unique borrows

Introduction to
Rust

Jacques-Henri
Jourdan

```
fn vec_inc(vec: &mut Vec<i32>) {  
    for el in vec {  
        *el += 1  
    }  
}  
  
fn main() {  
    let mut v = vec![18,5,7,2,9,27];  
    vec_inc(&mut v);  
    v.push(42); /* ... */  
}
```

Unique borrows are similar to shared borrows.

Differences:

- Allows mutation.
- Does not implement `Copy`.
- aka. **mutable borrow**, **mutable reference**.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Borrows and methods

Introduction to
Rust

Jacques-Henri
Jourdan

```
impl NumberOrNothing {
    fn print( self ) {
        match self {
            Nothing =>
                println!( "The number is: <nothing>" ),
            Number(n) =>
                println!( "The number is: {}", n ),
        };
    }
}
```

Call:

- copy of `self`
- full ownership transfer

Can we pass a borrow for `self`?

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Borrows and methods

Introduction to
Rust

Jacques-Henri
Jourdan

```
impl NumberOrNothing {
    fn print(&self) {
        match self {
            Nothing =>
                println!("The number is: <nothing>"),
            Number(n) =>
                println!("The number is: {}", n),
        };
    }
}
```

We can tell that `print` takes a borrow to `self`.

Same syntax at call site:

```
let n = Nothing;
n.print();
```

Borrowing is automatic at call-site for `self`.

Of course, this also works with `&mut`.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Pointers in data structures

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Imagine we have a vector of non-`Copy` elements. Say, `Vec<Vec<i32>>`.
How do we mutate the content?

One possibility:

```
fn set(i: i32, v: Vec<i32>, l: &mut Vec<Vec<i32>>)
```

The problem: cannot modify a `Vec<i32>` which is already in the vector!
We would need to move out the `Vec<i32>` and then move it in back. Not very efficient.

Pointers in data structures

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Imagine we have a vector of non-`Copy` elements. Say, `Vec<Vec<i32>>`.
How do we mutate the content?

Another possibility

```
fn get(i: i32, l: &mut Vec<Vec<i32>>) -> &mut Vec<i32>
```

but we said that borrows end when function terminate...

- We cannot use the return value

Borrows have **lifetimes**, noted `'a`, `'b`, `'c`...

- In general, for a borrow we write `&'a T` or `&mut 'a T`.
 - In some situations, this can be abbreviated to `&T` / `&mut T`.

Functions can be **lifetime polymorphic** to return inner pointer.

Example adapted from `Vec` library:

```
fn last_mut<'a, T>(v: &'a mut Vec<T>) -> Option<&'a mut T>
```

- Returns a borrow with the **same lifetime** than the parameter
- The calling function can update the content of the vector, as long as it does not access the vector at the same time
 - This preserves the ownership invariant (mutation XOR aliasing)
- The lifetime is **automatically inferred** when calling the function

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Lifetimes, example

Introduction to
Rust

Jacques-Henri
Jourdan

```
fn last_mut<'a, T>(v: &'a mut Vec<T>)
-> Option<&'a mut T>
{ ... }

fn f() {
    let mut v: Vec<i32> = vec![1,2,2];

    let opt = last_mut(&mut v);

    match opt {
        Some(last) => *last = 3,
        None => panic!()
    }
}
```

- When creating a borrow (with `&mut v`), Rust creates a **lifetime variable** `'a`.
- `v` is marked as **mutably borrowed** for `'a`.
 - `v` cannot be used as long as `'a` is alive.
- This lifetime variable goes through type inference.
⇒ variable `last` has type `&'a mut i32`
- We have only one constraint for `'a`: be alive when `last` is used
⇒ `'a` is inferred to be the orange zone.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Lifetimes, example

```
fn last_mut<'a, T>(v: &'a mut Vec<T>)
-> Option<&'a mut T>
{ ... }

fn f() {
    let mut v: Vec<i32> = vec![1,2,2];

    let opt = last_mut(&mut v);

    match opt {
        Some(last) => *last = 3,
        None => panic!()
    }

    v.push(4)
}
```

Let's add a mutation of `v` after the end of `'a`.

Rust checks the status of `v` at the call.
⇒ The lifetime has ended, `v` is available.
⇒ Call allowed.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Lifetimes, example

Introduction to
Rust

Jacques-Henri
Jourdan

```
fn last_mut<'a, T>(v: &'a mut Vec<T>)
-> Option<&'a mut T>
{ ... }

fn f() {
    let mut v: Vec<i32> = vec![1,2,2];

    let opt = last_mut(&mut v);

    v.push(42); // Danger!

    match opt {
        Some(last) => *last = 3,
        None => panic!()
    }
}
```

If we try to mutate `v` when it is still borrowed

Rust checks the status of `v` at the call.
The lifetime **has NOT ended**, the call is impossible.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Lifetimes, example

Introduction to
Rust

Jacques-Henri
Jourdan

```
fn last_mut<'a, T>(v: &'a mut Vec<T>)
    -> Option<&'a mut T>
{ ... }
```

```
fn f() {
    let mut v: Vec<u32> = Vec::new();

    let opt = last_mut(&mut v);
    v.push(42); // Danger!

    match opt {
        Some(last) => *last += 1,
        None => panic!("No element")
    }
}
```

```
error[E0499]: cannot borrow 'v' as mutable more than once at a time
--> src/lib.rs:10:3
   |
8 |     let opt = last_mut(&mut v);
   |                         ----- first mutable borrow occurs here
9 |
10|     v.push(42); // Danger!
   |             ----- second mutable borrow occurs here
11|
12|     match opt {
   |             --- first borrow later used here
```

If we try to mutate `v` when it is still borrowed

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Basic operations on borrows

Copying a shared borrow

(recall: `&T: Copy`):

```
let x = &1;
let y = x;
println!("{} {}", *x, *y);
```

Downcasting mutability:

```
let x = &mut 1;      // x: &'a mut i32
let y: &i32 = x;   // y: &'a i32
```

Splitting a borrow (shared or unique):

```
let x = &mut (1, 2); // x: &'a mut (i32, i32)
let (x1, x2) = x;   // x1: &'a mut i32  x2: &'b mut i32
```

Directly on a pair/record/enum:

```
let mut p = (42, 12);
let x = &mut p.0;
let y = &p.1; // Allowed even if p.0 is uniquely borrowed
let z = &p.1; // Allowed even if p.1 is already borrowed immutably
let t = p.1; // Idem
*x = 43; // x: &mut i32
println!("{} {} {}", *x, *y, *z); // y, z: &i32
```

Reborrowing

Introduction to
Rust

Jacques-Henri
Jourdan

```
fn last_mut<'a, T>(v: &'a mut Vec<T>)
    -> Option<&'a mut T>
{ ... }

let v = &mut vec![1, 2, 3];
match last_mut(&mut(*v)) { ... }
match last_mut(&mut(*v)) { ... }
```

Basic types

Aliasing control

Ownership in Rust

Borrowes and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Reborrowing creates a shorter borrow of the content of a borrow.
The old borrow is reactivated when the new borrow ends.

- Allows using a mutable borrow several times in a row.
- Happens implicitly in the vast majority of cases.
- The lifetime of the original borrow needs to be **longer** than the reborrowing lifetime.

Lifetime subtyping

Introduction to
Rust

Jacques-Henri
Jourdan

Reborrowing uses a notion of **lifetime inclusion**

- $\text{'a} \text{ outlives } \text{'b}$, written $\text{'a: } \text{'b}$.

This order relation: **lifetime subtyping**.

- The only source of subtyping in Rust.
- Constrained lifetime polymorphism:

```
fn f<'a, 'b: 'a>(...) -> ... { ... }
fn f<'a, 'b>(...) -> ... where 'b: 'a { ... }
```

- Exercise: what are the variances of `&`? `&mut`? `Box< >?` `Vec< >?`
 - With respect to the lifetime parameter?
 - With respect to the type parameter?

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Lifetime inference

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Most of the lifetime information is inferred by the compiler.

- We only need to annotate function types (and type declarations).

Key component of `rustc`: the **borrow checker**.

- Runs after traditional type-checking, and uses information from it.
- Uses variable liveness information from dataflow analysis.
- Tracks ownership (e.g., `Box<T>` value used only once) and infers lifetimes.
- Implemented on MIR, an intermediate representation in CFG form.

The borrow checker

Lifetime inference

The borrow checker interprets each lifetime `'a` as a set $\llbracket 'a \rrbracket$ containing:

- nodes of the CFG,
- elements `end('a)` for lifetimes `'a` generalized at the function's prototype.

Type-checking generates fresh lifetime variables and outlives constraints "`'a: 'b`:

- Examples: reborrowing, function calls, coercions, ...
- Interpreted as set inclusion $\llbracket 'b \rrbracket \subseteq \llbracket 'a \rrbracket$

Further constraints are added:

- $L \in \llbracket 'a \rrbracket$ when `'a` appears in the type of a live variable at CFG node `L`;
- $\text{end}('a) \in \llbracket 'a \rrbracket$ for any universal variable `'a`;
- $L \in \llbracket 'a \rrbracket$ for any universal variable `'a` and any CFG node `L`.

Borrow checker finds the smallest solution of the set of constraints (fixpoint algorithm).

The borrow checker

Checking the program

Once sets $\llbracket 'a \rrbracket$ are computed, it remains to check:

- that variables accesses are valid:
 - No read if not initialized or moved out.
 - No access if existing mutable loan with alive lifetime.
 - No write if existing immutable load with alive lifetime.
 - Question: when should borrowing be allowed?
- that outlives constraints at the function prototype are sufficient:
 - For any universally quantified lifetimes $'a$ and $'b$, if $\text{end}('a) \in \llbracket 'b \rrbracket$, then outlives constraints should imply $'b : 'a$.

The borrow checker

Checking the program

Once sets $\llbracket 'a \rrbracket$ are computed, it remains to check:

- that variables...
 - No reads after writes
 - No access to freed memory
 - No writes to freed memory
 - Questions about lifetimes
- that outlives constraints...
 - For any universally quantified lifetimes $'a$ and $'b$, if $\text{end}('a) \in \llbracket 'b \rrbracket$, then outlives constraints should imply $'b: 'a$.

Additional complications in `rustc`:

- Instead of variables: **places**
 - E.g., `rustc` treats `p.0` and `p.1` independently.
- Language features: closures...
- Efficient algorithm for fixpoint computation.
- ...

1 Basic types

Basic types

2 Aliasing control

- Ownership in Rust
- Borrows and lifetimes

Aliasing control

Ownership in Rust

Borrows and lifetimes

3 Generics & traits

- Iterators
- Closures
- Trait objects

Generics & traits

Iterators

Closures

Trait objects

Generics (aka. polymorphism)

What if we want to have a generic version of `vec_min`?

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;


```

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Generics (aka. polymorphism)

What if we want to have a generic version of `vec_min?`

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T>(vec: Vec<T>) -> Option<T> {
    let mut res = None;
    for el in vec {
        match res {
            None => { res = Some(el); },
            Some(n) => {
                let m = if n < el { n } else { el };
                res = Some(m);
            }
        }
    }
    return res
}
```

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Generics (aka. polymorphism)

What if we want to have a generic version of `vec_min`?

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T>(ve
let mut res =
for el in vec
match res {
    None => {
        Some(n) =>
            let m =
            res = So
    }
}
}

return res
}
```

Does this code compile?
Is there something missing?

Generics (aka. polymorphism)

What if we want to have a generic version of `vec_min`?

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T>(ve
let mut res =
for el in vec
match res {
    None => {
        Some(n) =>
            let m =
                res = So
}
}

return res
}
```

```
error[E0369]: binary operation '<' cannot be applied to type 'T'
--> src/lib.rs:13:22
|
13 |         let m = if n < el { n } else { el };
|             - ^ -- T
|             |
|             T
help: consider restricting type parameter 'T'
7  | fn vec_min<T: std::cmp::PartialOrd>(vec: Vec<T>) -> Option<T> {
|           ++++++
```

Generics (aka. polymorphism)

We have generics with trait bounds:

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T: std::cmp::PartialOrd>(vec: Vec<T>) -> Option<T> {
    let mut res = None;
    for el in vec {
        match res {
            None => { res = Some(el); },
            Some(n) => {
                let m = if n < el { n } else { el };
                res = Some(m);
            }
        }
    }
    return res
}
```

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Self: special type parameter for the `self` parameter of methods.

We can declare traits/type classes:

```
trait Trait<X1, X2, ...>: Trait0<...> {  
    type AssocTy;  
    fn method1(self, ...) -> Ty1;  
    fn method2(&self, ...) -> Ty2 {  
        // Default implementation  
    }  
    fn method3(x: i32, ...) -> Self;  
    ...  
}
```

```
class (C0 self ...) => C self a b ... where  
    data AssocTy self a b ... :: *;  
    method1 :: self -> ... -> ty1  
    method2 :: self -> ... -> ty2  
    method2 s ... =  
        -- Default implementation  
    method3 :: Int32 -> ... -> self  
    ...
```

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Self: special type parameter for the `self` parameter of methods.

We can write instances of them:

```
impl<X1, ...> Trait<Ty1, ....> for Ty
  where Ty2: Trait1, ... {

  type AssocTy = ...;
  fn method1(self, ...) -> Ty1 {
    ...
  }
  ...
}
```

```
instance (C1 ty2) => C ty1 ... where
  data AssocTy self ty1 ... = ...
  method1 s ... = ...
  ...
}
```

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Self: special type parameter for the `self` parameter of methods.

We can use them:

```
fn f<X>(...) -> ... where T: Trait {  
    // Same as: fn f<X: Trait>(...) -> ...  
    ...  
    x.method1(a, b, c)  
    ...  
}
```

```
f :: C x => ...  
f ... = ... method1 x a b c ...
```

```
fn vec_min<T: PartialOrd>(vec: Vec<T>) -> Option<T> {  
    ...  
    if n < el { n } else { el };  
    ...  
}  
  
trait PartialOrd { // Simplified  
    fn partial_cmp(&self, other: &Self) -> Option<Ordering>;  
    fn lt(&self, other: &Self) -> bool { ... }  
    ...  
}
```

```
fn vec_min<T: PartialOrd>(vec: Vec<T>) -> Option<T> {  
    ...  
    if n < el { n } else { el };  
    ...  
}  
  
trait PartialOrd { // Simplified  
    fn partial_cmp(&self, other: &Self) -> Option<Ordering>;  
    fn lt(&self, other: &Self) -> bool { ... }  
    ...  
}
```

Function `vec_min` is specialized for every type `T` with which it is used.
Hence, the call to `lt` is statically resolved (and can be inlined if necessary).

Comparison with Haskell's type classes

`Self` parameter

One of the trait parameter is implicit: `Self`.

- Just syntax, no real expressiveness change

Monomorphisation

Traits do not use indirect calls via dictionnaires like in Haskell.

Instead, **generic functions are specialized** for each use, and trait method calls are statically resolved.

This is a form of **zero cost abstraction**.

Higher-kinded parameters

Example: in `Monad M`, we have `M :: * -> *`.

In Rust, **only types** are allowed in trait parameters.

Coherence, orphan rules

There cannot be two conflicting instances in the same program.
This is **important for soundness** of some libraries.

Problem: how to enforce coherence when there are several compilation units?

Answer: “orphan rules”.

- Guarantees: no conflicting instances exist in two different compilation units.
- Roughly: one can declare an instance if trait or the implemented type is declared in the current trait.
- Details are complicated by generic types, parameterized traits, ...

Traits as a way of asserting properties

Some traits are used for expressing **properties of types**.

They do not have any method.

Examples:

- `Copy` expresses that a type can be copied without ownership issues.
- `Sized` expresses that the compiler knows the size of a type.
- `Send` and `Sync` express that a type is thread-safe (in some sense).

These traits have restrictions to which type can implement them:

- some of them are `unsafe`,
- others are treated specially by the compiler,
- some of them are automatically implemented (but we can opt-out).

An example of use of traits: Iterator and IntoIterator

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
trait Iterator {  
    type Item;  
    fn next(&mut self) -> Option<Self::Item>;  
    ...  
}  
  
trait IntoIterator {  
    type Item;  
    type IntoIter: Iterator<Item = Self::Item>;  
    fn into_iter(self) -> Self::IntoIter;  
}
```

Syntactic sugar:

```
for x in c {  
    ...  
} ⇒ let mut it = c.into_iter();  
loop {  
    match it.next() {  
        None => break,  
        Some(x) => {  
            ...  
        }  
    }  
}
```

An example of use of traits: Iterator and IntoIterator

```
trait Iterator {  
    type Item;  
    fn next(&&mut self) -> Option<Self::Item>;  
    ...  
}  
  
trait IntoIterator {  
    type Item;  
    type IntoIter;  
    fn into_iter(self) -> Self::IntoIter;  
}
```

Important to note: `Iterator` is not a type!

There is a different type for each implementation of `Iterator`.

This type describes the specialized state of the iterator. It is different depending on the data structure we are iterating over.

Syntactic sugar

Very different to e.g., the type `Seq.t` in OCaml.

```
for x in c {  
    ...  
}  
=====  
match it.next() {  
    None => break,  
    Some(x) => {  
        ...  
    }  
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

```
// The type behind the 0..N syntax
// (Specialized to i32 for simplicity)
struct Range {
    start: i32,
    end: i32,
}

impl Iterator for Range {
    type Item = i32;

    fn next(&mut self) -> Option<i32> {
        if self.start < self.end {
            let r = self.start; self.start += 1;
            return Some(r)
        } else {
            return None
        }
    }
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

The code written by the user:

```
for i in 0..N {  
    ...  
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

Removing syntactic sugar:

```
let r = Range{ start: 0, end: N };
let mut it = r; // IntoIterator is trivial for Range
loop {
    match it.next() {
        None => break,
        Some(i) => {
            ...
        }
    }
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

After inlining `next`:

```
let r = Range{ start: 0, end: N };
let mut it = r;
loop {
    let o =
        if it.start < end {
            let r = it.start; it.start += 1;
            Some(r)
        } else { None };
    match o {
        None => break,
        Some(i) => {
            ...
        }
    }
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

After inverting `if` and `match`, and simplifications:

```
let mut it = Range{ start: 0, end: N };
loop {
    if it.start < it.end {
        let i = it.start;
        it.start += 1;
        ...
    } else {
        break
    }
}
```

Very close to what one would expect from a C/C++ `for` loop!

It remains to:

- move `it.start += 1` to the end of the loop;
- use appropriate memory (stack or registers) for `it.start` and `it.end`.

Iterating over Vec<T>

There exists **three ways** to iterate over a vector:

```
for x in v { ... }      // x:T,      full ownership,      consumes v.  
for x in &mut v { ... } // x:&mut T, unique borrow of v, changes elements of v.  
for x in &v { ... }    // x:&T,      shared borrow of v, leaves v as-is.
```

Three different instances for `IntoIterator` (simplified):

```
impl<T> IntoIterator for Vec<T> {  
    type Item = T;  
    type IntoIter = IntoIter<T>;  
    fn into_iter(self) -> IntoIter<T> { ... }  
}  
  
impl<'a, T> IntoIterator for &'a Vec<T> {  
    type Item = &'a T;  
    type IntoIter = Iter<'a, T>;  
    fn into_iter(self) -> Iter<'a, T> { ... }  
}  
  
impl<'a, T> IntoIterator for &'a mut Vec<T> {  
    type Item = &'a mut T;  
    type IntoIter = IterMut<'a, T>;  
    fn into_iter(self) -> IterMut<'a, T> { ... }  
}
```

Iterating over Vec<T>

There exists **three ways** to iterate over a vector:

```
for x in v { ... }      // x:T,      full ownership,      consumes v.  
for x in &mut v { ... } // x:&mut T, unique borrow of v, changes elements of v.  
for x in &v { ... }    // x:&T,      shared borrow of v, leaves v as-is.
```

```
impl<'a, T> IntoIterator for &'a Vec<T> {  
    type Item = &'a T;  
    type IntoIter = Iter<'a, T>;  
    fn into_iter(self) -> Iter<'a, T> { ... }  
}
```

In borrow mode, the iterator (e.g., `Iter<'a, T>`) contains a borrow to the vector!

⇒ It depends on a lifetime `'a`.

⇒ `v` cannot be used as long as the iterator is alive.

⇒ No iterator invalidation!

Traditionally, in functional languages, closures are the runtime representation of values of the function type.

In Rust, two adjustments:

- The notion of **function type is too dynamic**.
Unsized \Rightarrow manipulated through a pointer.
Contains function pointer \Rightarrow indirect function calls (slow, inlining difficult).
 - In Rust, **functions are a trait**, not a type.
- Need to take into account ownership of captured variables.
 - **Three traits for closures:** `Fn`, `FnMut` and `FnOnce`.

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Closures

The FnOnce trait

```
trait FnOnce<Args> {  
    type Output;  
    fn call_once(self, args: Args) -> Self::Output;  
}
```

When the user writes `|a: &mut u64, b: i32| ...`, Rust generates during typing:

- an anonymous `struct` type, with one field for each captured variable,
- an instance of `FnOnce<(&mut u64, i32)>` for this anonymous `struct` type
 - `Output`: the return type of the function,
 - `call_once`: the body of the function (captured variables: via `self`),
- code which creates and initializes the `struct` of captured variables.

Notes:

- The `call_once` method gets full ownership of `self`.
 - The body of the closure can get **full ownership** of captured variables.
 - The closure can be called **only once!**
- `Args` is in fact a **tuple of parameters**.
- Notation for this trait: `FnOnce(&mut u64, i32) -> Option<u64>`

Closures

The FnOnce trait

```
trait FnOnce<Args> {  
    type Output;  
    fn call_once(self, args: Args) -> Self::Output;  
}
```

When the user writes `|a: &mut u64, b: i32| ...`, Rust generates during typing:

- an anonymous closure
 - an instance of `FnOnce<Args>`
 - `Output`
 - `call_once`
 - code which implements `FnOnce`
- This is a form of closure conversion, but during typing.
- Each closure gets a different type, but they all implement `FnOnce`.

Notes:

- The `call_once` method gets full ownership of `self`.
 - The body of the closure can get **full ownership** of captured variables.
 - The closure can be called **only once!**
- `Args` is in fact a **tuple of parameters**.
- Notation for this trait: `FnOnce(&mut u64, i32) -> Option<u64>`

Closures

The FnOnce trait, example

`std::iter::once(x)`: an iterator which returns `x`, only once.

```
fn once<A>(value: A) -> Once<A>
impl<A> Iterator for Once<A> { type Item = A; ... }
```

Closures

The FnOnce trait, example

`std::iter::once(x)`: an iterator which returns `x`, only once.

```
fn once<A>(value: A) -> Once<A>
impl<A> Iterator for Once<A> { type Item = A; ... }
```

A variant: `std::iter::once_with(f)`.

The same, but lazily, through a closure.

```
fn once_with<A, F>(gen: F) -> OnceWith<F>
    where F: FnOnce() -> A
impl<A, F: FnOnce() -> A> Iterator for OnceWith<F> {
    type Item = A;
    ...
}
```

Note: generic with respect to the closure type `F`!

Call example:

```
let it = std::iter::once_with(|| 42);
println!("{}", it.next());
```

Closures

Three traits: FnOnce, FnMut, Fn

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

```
trait FnOnce<Args> {
    type Output;
    fn call_once(self, args: Args) -> Self::Output;
}

trait FnMut<Args>: FnOnce<Args> { // Sub-trait
    fn call_mut(&mut self, args: Args) -> Self::Output;
}

trait Fn<Args>: FnMut<Args> { // Sub-trait
    fn call(&self, args: Args) -> Self::Output;
}
```

- Rust automatically generates as many instances as possible when using the `| ... | ...` syntax.
- Captured variables are moved/mutably borrowed only if necessary.

Examples of APIs with closures

Iterator combinators

There exists many **combinators for iterators**.

Examples:

```
trait Iterator {  
    ...  
    fn map<B, F>(self, f: F) -> Map<Self, F>  
        where F: FnMut(Self::Item) -> B { ... }  
    fn filter<P>(self, predicate: P) -> Filter<Self, P>  
        where P: FnMut(&Self::Item) -> bool { ... }  
    ...  
}  
  
impl<B, I: Iterator, F> Iterator for Map<I, F>  
where F: FnMut(I::Item) -> B,  
{ type Item = B; ... }  
  
impl<I: Iterator, P> Iterator for Filter<I, P>  
where P: FnMut(&I::Item) -> bool,  
{ type Item = I::Item; ... }
```

We can write, for example:

```
(0..N).filter(|i| i % 2 == 0).map(|i| i * i).sum()
```

Zero-cost abstraction with iterators and closures

Let's see how the following piece of code gets compiled:

```
(0..N).filter(|i| *i % 2 == 0).map(|i| i * i + a).sum()
```

We first expand closures...

Zero-cost abstraction with iterators and closures

After expanding closures:

```
struct C1 { a: i32 }
impl FnOnce<(i32,>) for C1 { type Output = i32; ... }
impl FnMut<(i32,>) for C1 {
    fn call_mut(&mut self, args:(i32,)) -> i32 { args.0 * args.0 + self.a }
}

struct C2 { /* No captured variable */ }
impl FnOnce<(&i32,>) for C2 { type Output = i32; ... }
impl FnMut<(&i32,>) for C2 {
    fn call_mut(&mut self, args:(&i32,)) -> bool { *i % 2 == 0 }
}

(0..N).filter(C1 { a: a }).map(C2 {}).sum()
```

We now inline `filter`, `map` and `sum`...

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Zero-cost abstraction with iterators and closures

After inlining `filter`, `map` and `sum`:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 {},
    iter:
        Filter {
            predicate: C1 { a: a },
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    match it.next() {
        None => break,
        Some(i) => r += i
    }
}
```

We now inline `<Map>::next...`

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

Zero-cost abstraction with iterators and closures

After inlining `<Map>::next`:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 {},
    iter:
        Filter {
            predicate: C1 { a: a },
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    let o = match it.iter.next() { None => None, Some(i) => Some(it.f.call_mut(i)) };
    match o {
        None => break,
        Some(i) => r += i
    }
}
```

Important: `call_mut` is a **direct call!**

We can merge the two `match` constructs, and inline `<C2>::call_mut`...

Zero-cost abstraction with iterators and closures

After inlining `<Filter>::next` and simplifications:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 {},
    iter:
        Filter {
            predicate: C1 { a: a },
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    match it.iter.next() {
        None => break,
        Some(i) => r += i*i+it.iter.predicate.a
    }
}
```

Similarly, we can inline `<Filter>::next` and `<C1>::call_mut...`

Zero-cost abstraction with iterators and closures

After inlining `<Filter>::next` and `<C1>::call_mut`:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 {},
    iter:
        Filter {
            predicate: C1 { a: a },
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    let o;
    loop { match it.iter.next() {
        None => { o = None; break },
        Some(i) => if i % 2 == 0 { o = Some(i); break }
    } }
    match o {
        None => break,
        Some(i) => r += i*i+it.predicate.a
    }
}
```

In the CFG graph, it is easy to see that the two `match` constructs can be merged...

Zero-cost abstraction with iterators and closures

We merged the two `match` constructs, and simplified:

```
/* Definition and instances of C1 and C2. */
...

let mut it = Map {
    f: C2 {},
    iter:
        Filter {
            predicate: C1 { a: a },
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    match it.iter.iter.next() {
        None => break,
        Some(i) => if i % 2 == 0 { r += i*i+it.iter.predicate.a }
    }
}
```

We inline `<Range>::next()` and expands the parts of `it` that stay unchanged...

Zero-cost abstraction with iterators and closures

... and we finally obtain:

```
let mut range_start = 0;
loop {
    if range_start < N {
        let i = range_start;
        range_start += 1;
        if i % 2 == 0 { r += i*i+a }
    } else { break }
}
```

We started with this concise and “abstract” piece of code:

```
(0..N).filter(|i| i % 2 == 0).map(|i| i * i + a).sum()
```

And ended-up with a well-optimized version!

No optimization step was difficult, any optimizing compiler performs these!

An example of zero-cost abstraction!

Trait objects

Sometimes, a true **function type** is really needed.

- Because too much specialization increases generated code size.
- Because sometimes the set of required specialization is unbounded (ex. code generated after CPS transform).

Solution: some traits can be turned into a type: **trait object**.

- Can be seen as a bounded existential type.
- Syntax: `dyn Trait`
- Trait objects types are **unsized**.
 - Size cannot be determined statically, depends on the implementing type.
 - Only used with (fat) pointers: `Box<dyn Trait>`, `&dyn Trait`, ...
- Method calls are performed indirectly, through a vtable.
 - Some methods are incompatible: generics, using `Self` without a pointer...
⇒ only **object safe** traits are allowed in a trait object.

Trait objects

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

What type would you use for continuations in CPS?

Trait objects

Introduction to
Rust

Jacques-Henri
Jourdan

Basic types

Aliasing control

Ownership in Rust

Borrows and lifetimes

Generics & traits

Iterators

Closures

Trait objects

What type would you use for continuations in CPS?

We could use `Box<dyn FnOnce(...) -> ...>`.

Conclusion: Rust uses the **same abstraction mechanism** for both static dispatch (i.e., templates in C++) and dynamic dispatch (i.e., virtual methods in C++).