

Rust: When the Aliasing Discipline Is Too Strong

Jacques-Henri Jourdan

January 31st, 2024

Abstract from last week

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Last week, we saw many features of the Rust programming language:

- algebraic data types, records;
- various kinds of pointers: `Box<T>`, `&mut T`, `&T`;
- how Rust controls ownership and aliasing, the concept of lifetime;
- traits: Rust's type classes
 - (trait examples: iterators, closures);
- why all that gives control, performance **and** safety
 - (the notion of zero-cost abstraction).

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

`Cell<T>`

`RefCell<T>`

`Rc<T>`

This week

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

How to escape the ownership discipline when needed?

- unsafe code behind safe abstractions,
- interior mutability.

1 Unsafe Rust

- Unsafe and aliasing
- Safe abstractions

Unsafe Rust

Unsafe and aliasing
Safe abstractions

Interior mutability

Cell<T>
RefCell<T>
Rc<T>

2 Interior mutability

- Cell<T>
- RefCell<T>
- Rc<T>

Unsafe

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Despite zero-cost abstractions, the requirement of safety has some costs:

- no loop in memory graph,
- shared ownership: restricted to borrows (statically known lifetime),
- linking to external libraries: not always follow ownership discipline,
- bounds checks for `Vec` accesses.

One can use **unsafe Rust** to avoid these costs.

- A set of features that extend Rust.
- Only available in `unsafe` blocks or `unsafe` functions.
- No guarantee of safety:
 - one should **encapsulate unsafety behind safe abstractions**,
 - or mark functions with unsafe behavior as `unsafe`.

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

`Cell<T>`

`RefCell<T>`

`Rc<T>`

Unsafe features

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing
Safe abstractions

Interior mutability
Cell<T>
RefCell<T>
Rc<T>

- Dereference **raw pointers**.
- Call unsafe functions (e.g., accessing `Vec<T>` without bounds checks).
- Implement **unsafe** traits (e.g., `Send` and `Sync`, see next week).

And things we won't discuss:

- Mutate global variables (called `static` variables in Rust)
 - (because that can lead to data races).
- Access `union` types.

Undefined behaviors

Using these features is **unsafe**.

⇒ the compiled program can crash even if type-checked!

It is important to know when a program triggers **undefined behavior**.

This is sometimes **very subtle**.

Rust provides an experimental **reference interpreter**, called **Miri**, which detects **some** undefined behavior.

- Can be used on concrete code for testing.

Still, writing unsafe code should be **reserved to experts**. There is a book dedicated to writing unsafe code: **Rustonomicon**.

So I will not teach this here. Instead, we will see why this is subtle.

Unsafe and aliasing

Common use of unsafe code: weaken aliasing restrictions.

Raw pointers `*mut T` and `*const T`:

- have no statically checked aliasing restriction,
- can coerce from/to borrows (both shared and unique),
- can easily be used to break any aliasing policy.

Unsafe and aliasing

Common use of unsafe code: weaken aliasing restrictions.

Raw pointers `*mut T` and `*const T`:

- have no statically checked aliasing restriction,
- can coerce from/to borrows (both shared and unique),
- can easily be used to break any aliasing policy.

But the compiler may assume known aliasing properties on borrows to perform some optimizations:

```
fn test_noalias(x: &mut i32, y: &mut i32) -> i32 {  
    // x, y cannot alias: they are unique borrows  
    *x = 42;  
    *y = 37;  
    return *x; // must return 42 -- can be optimized  
}
```

⇒ The programmer should take care of not breaking these aliasing guarantees using raw pointers!

Unsafe and aliasing

Common use of unsafe code: weaken aliasing restrictions.

Raw pointers `*mut T` and `*const T`:

- have no statically checked aliasing restriction,
- can coerce from/to borrows (both shared and unique),
- can easily be used to break any aliasing policy.

But the compiler may assume known aliasing properties on borrows to perform some optimizations:

```
fn test_unique(x: &mut i32) -> i32 {  
    *x = 42;  
    // unknown_function cannot have an alias to x  
    unknown_function();  
    return *x; // must return 42 -- can be optimized  
}
```

⇒ The programmer should take care of not breaking these aliasing guarantees using raw pointers!

Unsafe and aliasing

Common use of unsafe code: weaken aliasing restrictions.

Raw pointers `*mut T` and `*const T`:

- have no statically checked aliasing restriction,
- can coerce from/to borrows (both shared and unique),
- can easily be used to break any aliasing policy.

But the compiler may assume known aliasing properties on borrows to perform some optimizations:

```
fn test_shared(x: &i32) -> i32 {  
    let y = *x;  
    // unknown_function cannot have a mutable alias to x  
    unknown_function();  
    return *x + y; // can be optimized to 2*y  
}
```

⇒ The programmer should take care of not breaking these aliasing guarantees using raw pointers!

Undefined behavior and aliasing

Some rules are needed to tell what one can do with raw pointers.

These rules must be a balance between:

- flexibility for the programmer of unsafe code;
- allowing optimizations.

Choosing these rules is still an [open problem](#).

The Miri interpreter implements two sets of rules called [Stacked Borrows](#) and [Tree Borrows](#):

- experimental and imperfect,
- but executable on concrete tests.

Undefined behavior and aliasing

Some rules are needed to tell what one can do with raw pointers.

These rules must be a balance between:

- flexibility for users
- allowing operations

If you write unsafe code, you need to follow rules like these.

Choosing these

The Miri interpreter

Borrows:

I told you, writing correct unsafe code is subtle...

- experimental and imperfect,
- but executable on concrete tests.

Safe abstractions

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

How to benefit from the power of unsafe code without paying the cost?

Use **libraries** written using unsafe features but with safe interfaces.

Example: `Vec<T>`, resizable arrays:

- fully written in Rust with unsafe features,
- yet, most of the functions exposed by `Vec<T>` are safe!

Example of a safe abstraction: a queue based on a linked list

Let's say we would like to implement a FIFO queue using a singly linked list.

We need a pointer both at the beginning (for `pop`) and at the end of the list (for `push`).

Aliasing rules are violated, we need unsafe code.

Example of safe abstraction: Queue<T>

```
mod queue {  
    pub struct Queue<T> {  
        head: *mut Node<T>,  
        tail: *mut Node<T>  
    }  
    struct Node<T> {  
        elem: T,  
        next: *mut Node<T>,  
    }  
    ...  
}  
use queue::*;

fn (q: Queue<i32> /* Allowed */) {  
    ...  
    q.head /* Error */  
    ...  
    Queue { ... } /* Error */  
    ...  
    let x : Node<i32> /* Error */ = ... ;  
    ...  
}
```

We use **modules** as an encapsulation mechanism.

Some elements of the modules are marked with **pub**, they are public.

The other elements are private.

Fields of **struct** are also either public or private.

- Queue has no public field: **abstract type!**

1 Unsafe Rust

- Unsafe and aliasing
- Safe abstractions

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

2 Interior mutability

- Cell<T>
- RefCell<T>
- Rc<T>

Working around aliasing rules

Can we do better than `unsafe`?

On the one hand, Rust aliasing rules are strict;
on the other hand, `unsafe` code seems too subtle to write...

Can we do better?

Working around aliasing rules

Can we do better than `unsafe`?

On the one hand, Rust aliasing rules are strict;
on the other hand, `unsafe` code seems too subtle to write...

Can we do better?

We can use **interior mutability**:

- libraries that relax aliasing rules, **safely**,
- written with `unsafe` code, but safely encapsulated!
- Common feature: updating memory using a shared borrow, with appropriate restrictions.
 - (Uses special annotation to disable some optimizations.)

Any idea of an API with interior mutability?

```
pub struct Cell<T> { ... }

impl<T> Cell<T> {
    pub fn new(value: T) -> Cell<T> { ... }
    pub fn into_inner(self) -> T { ... }
    pub fn set(&self, val: T) { ... }
    pub fn replace(&self, val: T) -> T { ... }
    pub fn get(&self) -> T where T : Copy { ... }
}
```

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Informally, why is this safe?

```
pub struct Cell<T> { ... }

impl<T> Cell<T> {
    pub fn new(value: T) -> Cell<T> { ... }
    pub fn into_inner(self) -> T { ... }
    pub fn set(&self, val: T) { ... }
    pub fn replace(&self, val: T) -> T { ... }
    pub fn get(&self) -> T where T : Copy { ... }
}
```

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Informally, why is this safe?

From `&Cell<T>`, we can never get a (shared or mutable) borrow of the content. Hence, invariants on borrows of `T` cannot be violated.

We can only exchange values of type `T` or get a copy, but no internal borrow.

```
pub struct Cell<T> { ... }

impl<T> Cell<T> {
    pub fn new(value: T) -> Cell<T> { ... }
    pub fn into_inner(self) -> T { ... }
    pub fn set(&self, val: T) { ... }
    pub fn replace(&self, val: T) -> T { ... }
    pub fn get(&self) -> T where T : Copy { ... }
}
```

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Informally, why is this safe?

From `&Cell<T>`, we can never get a (shared or mutable) borrow of the content. Hence, invariants on borrows of `T` cannot be violated.

We can only exchange values of type `T` or get a copy, but no internal borrow.

And what if we **do want** an internal borrow?

RefCell<T> API(1/2)

RefCell, RefMut

```
pub struct RefCell<T> { ... }  
pub struct RefMut<'b, T> where T: 'b { ... }  
  
impl<T> RefCell<T> {  
    pub fn new(value: T) -> RefCell<T> { ... }  
    pub fn into_inner(self) -> T { ... }  
  
    /* Checks there is no borrow. Marks as uniquely borrowed. */  
    pub fn borrow_mut<'a>(&'a self) -> RefMut<'a, T> { ... }  
}  
  
/* This DerefMut instance means RefMut<'b, T> can be used as &'b mut T*/  
impl<'b, T> DerefMut for RefMut<'b, T> {  
    fn deref_mut<'a>(&'a mut self) -> &'a mut T /* where 'b: 'a */ { ... }  
}  
/* This Deref instance means RefMut<'b, T> can be used as &'b T */  
impl<'b, T> Deref for RefMut<'b, T> {  
    type Target = T  
    fn deref<'a>(&'a self) -> &'a T /* where 'b: 'a */ { ... }  
}  
  
/* Destructor. */  
impl<'a, T> Drop for RefMut<'a, T> {  
    /* Mark RefCell as not borrowed. */  
    fn drop(&mut self) { ... }  
}
```

RefCell<T> Example

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing
Safe abstractions

Interior mutability

Cell<T>
RefCell<T>
Rc<T>

```
fn use_refcell(x : &RefCell<Vec<i32>>) {  
    let mut v: RefMut<'_, Vec<i32>> = x.borrow_mut();  
    v.push(42); // v can be used just like a unique borrow  
  
    /* Panics: there already is a unique borrow. */  
    /* let v2 = x.borrow_mut().push(16); */  
  
    /* Implicit : v.drop(); */  
}  
  
/* The RefMut is dropped, I can create another one: */  
println!("{}", x.borrow_mut()[0]);  
}
```

RefCell<T> API (2/2)

Ref

```
...
pub struct Ref<'b, T> where T: 'b { ... }

impl<T> RefCell<T> {
    ...
    /* Checks there is no unique borrow. Increments borrow count. */
    pub fn borrow<'a>(&'a self) -> Ref<'a, T> { ... }
}

/* This Deref instance means Ref<'b, T> can be used as &'b T */
impl<'b, T> Deref for Ref<'b, T> {
    type Target = T
    fn deref<'a>(&'a self) -> &'a T /* where 'b: 'a */ { ... }
}

/* Destructor. */
impl<'a, T> Drop for Ref<'a, T> {
    /* Decrement borrow count. */
    fn drop(&mut self) { ... }
}
```

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing
Safe abstractions

Interior mutability

Cell<T>
RefCell<T>
Rc<T>

Soundness

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Why is RefCell sound?

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Soundness

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Why is RefCell sound?

The aliasing rule (aliasing XOR mutation) is enforced dynamically, through an internal counter.

We can see RefCell as a non-concurrent reader/writer lock.

Subtle API question regarding lifetimes

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

In the standard library:

```
impl<'b, T> RefMut<'b, T> {
    pub fn map<U, F>(orig: RefMut<'b, T>, f: F) -> RefMut<'b, U>
        where F: FnOnce(&mut T) -> &mut U
    { ... }
}

impl<'b, T> Ref<'b, T> {
    pub fn map<U, F>(orig: Ref<'b, T>, f: F) -> Ref<'b, U>
        where F: FnOnce(&T) -> &U
    { ... }
}
```

It can be used e.g., to transform a `RefMut<'b, T>` to a `RefMut` to one of the fields of `T`.

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Subtle API question regarding lifetimes

In the standard library:

```
impl<'b, T> RefMut<'b, T> {  
    pub fn map<U, F>(orig: RefMut<'b, T>, f: F) -> RefMut<'b, U>  
        where F: FnOnce(&mut T) -> &mut U  
    { ... }  
}  
  
impl<'b, T> Ref<'b, T> {  
    pub fn map<U, F>(orig: Ref<'b, T>, f: F) -> Ref<'b, U>  
        where F: FnOnce(&T) -> &U  
    { ... }  
}
```

It can be used e.g., to transform a `RefMut<'b, T>` to a `RefMut` to one of the fields of `T`.

Exercise: what are the lifetimes of borrows  used in closures?
Give (counter-)examples.

```
struct Rc<T> { ... }

impl<T> Rc<T> {
    pub fn new(value: T) -> Rc<T> { ... }
}

/* This Deref instance means Rc<T> can be used as &T */
impl<T> Deref for Rc<T> {
    type Target = T
    fn deref<'a>(&'a self) -> &'a T { ... }
}

impl<T> Clone for Rc<T> {
    /* Copy the pointer, and increment the reference count. */
    fn clone(&self) -> Rc<T> { ... }
}

impl<T> Drop for Rc<T> {
    /* Drop the pointer, decrement the reference count, and recursively
       drop+deallocate if count is 0. */
    fn drop(&mut self) { ... }
}
```

Rc<T>

A pointer to T, with reference counting

```
struct Rc<T> { ... }

impl<T> Rc<T> {
    pub fn new(value: T) -> Rc<T> { ... }
}

/* This Deref implements Deref for Target */
impl<T> Deref for Target {
    type Target = T;
    fn deref<'a>(&'a self) -> &'a T
}

impl<T> Clone for Rc<T> {
    /* Copy the value */
    fn clone(&self) -> Rc<T> { ... }
}

impl<T> Drop for Rc<T> {
    /* Drop the pointer, decrement the reference count, and recursively
       drop+deallocate if count is 0. */
    fn drop(&mut self) { ... }
}
```

This is typically used for implementing data structures with sharing.
Example: purely functional maps, BDDs...

Why do I say this is interior mutability?

Rc<T>

A pointer to T, with reference counting

```
struct Rc<T> { ... }

impl<T> Rc<T> {
    pub fn new(value: T) -> Rc<T> { ... }
}

/* This Deref implements Deref for Target */
impl<T> Deref for Target {
    type Target = T;
    fn deref<'a>(&'a self) -> &'a T
}

impl<T> Clone for Rc<T> {
    /* Copy the value */
    fn clone(&self) -> Rc<T> { ... }
}

impl<T> Drop for Rc<T> {
    /* Drop the pointer, decrement the reference count, and recursively
       drop+deallocate if count is 0. */
    fn drop(&mut self) { ... }
}
```

This is typically used for implementing **data structures with sharing**.
Example: purely functional maps, BDDs...

Interior mutability is limited to the reference count.

Getting mutable references

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

A priori, an `Rc<T>` can be aliased, so we don't have a `DerefMut` instance.

But we have:

```
impl<T> Rc<T> {
    /* Checks the count is equal to 1. */
    pub fn get_mut(this: &mut Rc<T>) -> Option<&mut T> { ... }

    /* Clone-on-write: clone the content to a fresh location if the count is not 1. */
    pub fn make_mut(this: &mut Rc<T>) -> &mut T
        where T : Clone { ... }
}
```

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Of course, this prevents mutation and aliasing.

How would you get a functionality close to an OCaml's ref type?

Getting mutable references

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

A priori, an `Rc<T>` can be aliased, so we don't have a `DerefMut` instance.

But we have:

```
impl<T> Rc<T> {
    /* Checks the count is equal to 1. */
    pub fn get_mut(this: &mut Rc<T>) -> Option<&mut T> { ... }

    /* Clone-on-write: clone the content to a fresh location if the count is not 1. */
    pub fn make_mut(this: &mut Rc<T>) -> &mut T
        where T : Clone { ... }
}
```

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Of course, this prevents mutation and aliasing.

How would you get a functionality close to an OCaml's `ref` type?

Answer: `Rc<RefCell<T>>`.

A note on performances

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

Reference counting is sometimes considered slow.

This is because it usually requires a lot of updates to the reference counts (ex. parameter passing, assignments to a variable...)

In Rust, we can mix reference counting and borrowing:

- Use borrows when doing a read-only traversal of a data structure.
- Increment the count only when creating a new long-lived reference.

This gives better control, and better performances.

A question for next course

Rust: When the
Aliasing Discipline
Is Too Strong

Jacques-Henri
Jourdan

Unsafe Rust

Unsafe and aliasing

Safe abstractions

Interior mutability

Cell<T>

RefCell<T>

Rc<T>

The next course will focus on support for parallelism (i.e., multicore programs).

What do you think is needed in Rust to guarantee the safety of
multicore Rust programs?