

Rust: Generics and Traits

Jacques-Henri Jourdan

February 4th, 2025

Do not forget your project.

Deadline: Friday, February, 28th.

Send it to the following email address:

jacques-henri.jourdan@universite-paris-saclay.fr

(My email address @cnrs.fr has bad “security” filters which removes your attachment.)

Do not hesitate to ask questions **early**.

Abstract from last time (10/12/2024)

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

Last time, we saw many features of the Rust programming language:

- algebraic data types, records;
- various kinds of pointers: `Box<T>`, `&mut T`, `&T`;
- how Rust controls ownership and aliasing, the concept of lifetime;
- Unsafe blocks/functions: workaround strong static type-checking constraints,
 - but correct unsafe code: very subtle.
- Encapsulation: clients can safely use libraries written with unsafe code.
- Interior mutability (the ability to mutate through shared borrows): a typical example of well-encapsulated unsafe code.

Example from last week

Rust: Generics
and Traits

Jacques-Henri
Jourdan

```
enum NumberOrNothing {
    Number(i32), Nothing
}
use NumberOrNothing::*;

fn vec_min(vec: Vec<i32>) -> {
    let mut res = Nothing;
    for el in vec {
        match res {
            Nothing => {
                res = Number(el);
            },
            Number(n) => {
                let m = if n < el { n } else { el };
                res = Number(m);
            }
        }
    }
    return res
}
```

Iterators
Closures
Trait objects

Generics (aka. polymorphism)

Rust: Generics
and Traits

Jacques-Henri
Jourdan

What if we want to have a generic version of `vec_min`?

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;


```

Iterators
Closures
Trait objects

Generics (aka. polymorphism)

Rust: Generics
and Traits

Jacques-Henri
Jourdan

What if we want to have a generic version of `vec_min?`

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T>(vec: Vec<T>) -> Option<T> {
    let mut res = None;
    for el in vec {
        match res {
            None => { res = Some(el); },
            Some(n) => {
                let m = if n < el { n } else { el };
                res = Some(m);
            }
        }
    }
    return res
}
```

Iterators
Closures
Trait objects

Generics (aka. polymorphism)

What if we want to have a generic version of `vec_min`?

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T>(ve
let mut res =
for el in vec
match res {
    None => {
        Some(n) =>
            let m =
            res = So
    }
}
}

return res
}
```

Does this code compile?
Is there something missing?

Generics (aka. polymorphism)

What if we want to have a generic version of `vec_min`?

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T>(ve
let mut res =
for el in vec
match res {
    None => {
        Some(n) =>
            let m =
                res = So
    }
}
}

return res
}
```

```
error[E0369]: binary operation '<' cannot be applied to type 'T'
--> src/lib.rs:13:22
|
13 |         let m = if n < el { n } else { el };
|             - ^ -- T
|             |
|             T
help: consider restricting type parameter 'T'
7  | fn vec_min<T: std::cmp::PartialOrd>(vec: Vec<T>) -> Option<T> {
|               ++++++
```

Generics (aka. polymorphism)

Rust: Generics
and Traits

Jacques-Henri
Jourdan

We have generics with trait bounds:

```
enum Option<T> { // Already in stdlib
    None,
    Some(T),
}
use Option::*;

fn vec_min<T: std::cmp::PartialOrd>(vec: Vec<T>) -> Option<T> {
    let mut res = None;
    for el in vec {
        match res {
            None => { res = Some(el); },
            Some(n) => {
                let m = if n < el { n } else { el };
                res = Some(m);
            }
        }
    }
    return res
}
```

Iterators
Closures
Trait objects

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Self: special type parameter for the `self` parameter of methods.

We can declare traits/type classes:

```
trait Trait<X1, X2, ...>: Trait0<...> {  
    type AssocTy;  
    fn method1(self, ...) -> Ty1;  
    fn method2(&self, ...) -> Ty2 {  
        // Default implementation  
    }  
    fn method3(x: i32, ...) -> Self;  
    ...  
}
```

```
class (C0 self ...) => C self a b ... where  
    data AssocTy self a b ... :: *;  
    method1 :: self -> ... -> ty1  
    method2 :: self -> ... -> ty2  
    method2 s ... =  
        -- Default implementation  
    method3 :: Int32 -> ... -> self  
    ...
```

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Self: special type parameter for the `self` parameter of methods.

We can write instances of them:

```
impl<X1, ...> Trait<Ty1, ....> for Ty
  where Ty2: Trait1, ... {

  type AssocTy = ...;
  fn method1(self, ...) -> Ty1 {
    ...
  }
  ...
}
```

```
instance (C1 ty2) => C ty1 ... where
  data AssocTy self ty1 ... = ...
  method1 s ... = ...
  ...
}
```

Traits (Rust's type classes)

Rust's type traits resemble Haskell's type classes:

- A mechanism for **overloading**:
 - give different meaning to the same function name or symbol, depending on parameters and return types.
- A mechanism for **abstraction**:
 - the same function works for different implementations of some interface.

Self: special type parameter for the `self` parameter of methods.

We can use them:

```
fn f<X>(...) -> ... where X: Trait {  
    // Same as: fn f<X: Trait>(...) -> ...  
    ...  
    x.method1(a, b, c)  
    ...  
}
```

```
f :: C x => ...  
f ... = ... method1 x a b c ...
```

Function vec_min requires T: PartialOrd

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

```
fn vec_min<T: PartialOrd>(vec: Vec<T>) -> Option<T> {  
    ...  
    if n < el { n } else { el };  
    ...  
}  
  
trait PartialOrd { // Simplified  
    fn partial_cmp(&self, other: &Self) -> Option<Ordering>;  
    fn lt(&self, other: &Self) -> bool { ... }  
    ...  
}
```

```
fn vec_min<T: PartialOrd>(vec: Vec<T>) -> Option<T> {  
    ...  
    if n < el { n } else { el };  
    ...  
}  
  
trait PartialOrd { // Simplified  
    fn partial_cmp(&self, other: &Self) -> Option<Ordering>;  
    fn lt(&self, other: &Self) -> bool { ... }  
    ...  
}
```

Function `vec_min` is specialized for every type `T` with which it is used.
Hence, the call to `lt` is statically resolved (and can be inlined if necessary).

Comparison with Haskell's type classes

`Self` parameter

One of the trait parameter is implicit: `Self`.

- Just syntax, no real expressiveness change

Monomorphisation

Traits do not use indirect calls via dictionnaires like in Haskell.

Instead, **generic functions are specialized** for each use, and trait method calls are statically resolved.

This is a form of **zero cost abstraction**.

Higher-kinded parameters

Example: in `Monad M`, we have `M :: * -> *`.

In Rust, **only types** are allowed in trait parameters.

Coherence, orphan rules

There cannot be two conflicting instances in the same program.
This is **important for soundness** of some libraries.

Problem: how to enforce coherence when there are several compilation units?

Answer: “orphan rules”.

- Guarantees: no conflicting instances exist in two different compilation units.
- Roughly: one can declare an instance if trait or the implemented type is declared in the current trait.
- Details are complicated by generic types, parameterized traits, ...

Traits as a way of asserting properties

Some traits are used for expressing **properties of types**.

They do not have any method.

Examples:

- `Copy` expresses that a type can be copied without ownership issues.
- `Sized` expresses that the compiler knows the size of a type.
- `Send` and `Sync` express that a type is thread-safe (in some sense).

These traits have restrictions to which type can implement them:

- some of them are `unsafe`,
- others are treated specially by the compiler,
- some of them are automatically implemented (but we can opt-out).

An example of use of traits: Iterator and IntoIterator

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

```
trait Iterator {
    type Item;
    fn next(&mut self) -> Option<Self::Item>;
    ...
}

trait IntoIterator {
    type Item;
    type IntoIter: Iterator<Item = Self::Item>;
    fn into_iter(self) -> Self::IntoIter;
}
```

Syntactic sugar:

```
for x in c {  
    ...  
}  
=====  
let mut it = c.into_iter();  
loop {  
    match it.next() {  
        None => break,  
        Some(x) => {  
            ...  
        }  
    }  
}
```

An example of use of traits: Iterator and IntoIterator

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

```
trait Iterator {  
    type Item;  
    fn next(&mut self) -> Option<Self::Item>;  
    ...  
}  
  
trait IntoIterator {  
    type Item;  
    type IntoIter;  
    fn into_iter(self) -> IntoIter;  
}
```

Important to note: `Iterator` is not a type!

There is a different type for each implementation of `Iterator`.

This type describes the specialized state of the iterator. It is different depending on the data structure we are iterating over.

Syntactic sugar

Very different to e.g., the type `Seq.t` in OCaml.

```
for x in c {  
    ...  
}  $\Rightarrow$  match it.next() {  
    None => break,  
    Some(x) => {  
        ...  
    }  
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

Ranges

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

```
// The type behind the 0..N syntax
// (Specialized to i32 for simplicity)
struct Range {
    start: i32,
    end: i32,
}

impl Iterator for Range {
    type Item = i32;

    fn next(&mut self) -> Option<i32> {
        if self.start < self.end {
            let r = self.start; self.start += 1;
            return Some(r)
        } else {
            return None
        }
    }
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

The code written by the user:

```
for i in 0..N {  
    ...  
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

Removing syntactic sugar:

```
let r = Range{ start: 0, end: N };
let mut it = r; // IntoIterator is trivial for Range
loop {
    match it.next() {
        None => break,
        Some(i) => {
            ...
        }
    }
}
```

Ranges

Rust: Generics
and Traits

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

Jacques-Henri
Jourdan

Iterators

Closures

Trait objects

After inlining `next`:

```
let r = Range{ start: 0, end: N };
let mut it = r;
loop {
    let o =
        if it.start < end {
            let r = it.start; it.start += 1;
            Some(r)
        } else { None };
    match o {
        None => break,
        Some(i) => {
            ...
        }
    }
}
```

Ranges

- Ranges such as `0..10` implement `IntoIterator` and `Iterator`.
- `for i in 0..N { ... }` is easily optimized.
 - This is actually the usual way of writing `for` loops.

After inverting `if` and `match`, and simplifications:

```
let mut it = Range{ start: 0, end: N };
loop {
    if it.start < it.end {
        let i = it.start;
        it.start += 1;
        ...
    } else {
        break
    }
}
```

Very close to what one would expect from a C/C++ `for` loop!

It remains to:

- move `it.start += 1` to the end of the loop;
- use appropriate memory (stack or registers) for `it.start` and `it.end`.

Iterating over Vec<T>

Rust: Generics
and Traits

Jacques-Henri
Jourdan

There exists **three ways** to iterate over a vector:

```
for x in v { ... }      // x:T,      full ownership,      consumes v.  
for x in &mut v { ... } // x:&mut T, unique borrow of v, changes elements of v.  
for x in &v { ... }    // x:&T,      shared borrow of v, leaves v as-is.
```

Three different instances for **IntoIterator** (simplified):

```
impl<T> IntoIterator for Vec<T> {  
    type Item = T;  
    type IntoIter = IntoIter<T>;  
    fn into_iter(self) -> IntoIter<T> { ... }  
}  
  
impl<'a, T> IntoIterator for &'a Vec<T> {  
    type Item = &'a T;  
    type IntoIter = Iter<'a, T>;  
    fn into_iter(self) -> Iter<'a, T> { ... }  
}  
  
impl<'a, T> IntoIterator for &'a mut Vec<T> {  
    type Item = &'a mut T;  
    type IntoIter = IterMut<'a, T>;  
    fn into_iter(self) -> IterMut<'a, T> { ... }  
}
```

Iterators
Closures
Trait objects

Iterating over Vec<T>

There exists **three ways** to iterate over a vector:

```
for x in v { ... }      // x:T,      full ownership,      consumes v.  
for x in &mut v { ... } // x:&mut T, unique borrow of v, changes elements of v.  
for x in &v { ... }    // x:&T,      shared borrow of v, leaves v as-is.
```

```
impl<'a, T> IntoIterator for &'a Vec<T> {  
    type Item = &'a T;  
    type IntoIter = Iter<'a, T>;  
    fn into_iter(self) -> Iter<'a, T> { ... }  
}
```

In borrow mode, the iterator (e.g., `Iter<'a, T>`) contains a borrow to the vector!

⇒ It depends on a lifetime `'a`.

⇒ `v` cannot be used as long as the iterator is alive.

⇒ No iterator invalidation!

Traditionally, in functional languages, closures are the runtime representation of values of the function type.

In Rust, two adjustments:

- The notion of **function type is too dynamic**.
Unsized \Rightarrow manipulated through a pointer.
Contains function pointer \Rightarrow indirect function calls (slow, inlining difficult).
 - In Rust, **functions are a trait**, not a type.
- Need to take into account ownership of captured variables.
 - **Three traits for closures:** `Fn`, `FnMut` and `FnOnce`.

Closures

The FnOnce trait

```
trait FnOnce<Args> {  
    type Output;  
    fn call_once(self, args: Args) -> Self::Output;  
}
```

When the user writes `|a: &mut u64, b: i32| ...`, Rust generates **during typing**:

- an anonymous `struct` type, with one field for each captured variable,
- an instance of `FnOnce<(&mut u64, i32)>` for this anonymous `struct` type
 - `Output`: the return type of the function,
 - `call_once`: the body of the function (captured variables: via `self`),
- code which creates and initializes the `struct` of captured variables.

Notes:

- The `call_once` method gets full ownership of `self`.
 - The body of the closure can get **full ownership** of captured variables.
 - The closure can be called **only once!**
- `Args` is in fact a **tuple of parameters**.
- Notation for this trait: `FnOnce(&mut u64, i32) -> Option<u64>`

Closures

The FnOnce trait

```
trait FnOnce<Args> {  
    type Output;  
    fn call_once(self, args: Args) -> Self::Output;  
}
```

When the user writes `|a: &mut u64, b: i32| ...`, Rust generates **during typing**:

- an anonymous closure
 - an instance of `FnOnce<Args>`
 - `Output`
 - `call_once`
 - code which implements `FnOnce<Args>`
- This is a form of closure conversion, but during typing.
- Each closure gets a different type, but they all implement `FnOnce`.

Notes:

- The `call_once` method gets full ownership of `self`.
 - The body of the closure can get **full ownership** of captured variables.
 - The closure can be called **only once!**
- `Args` is in fact a **tuple of parameters**.
- Notation for this trait: `FnOnce(&mut u64, i32) -> Option<u64>`

Closures

The FnOnce trait, example

`std::iter::once(x)`: an iterator which returns `x`, only once.

```
fn once<A>(value: A) -> Once<A>
impl<A> Iterator for Once<A> { type Item = A; ... }
```

Closures

The FnOnce trait, example

`std::iter::once(x)`: an iterator which returns `x`, **only once**.

```
fn once<A>(value: A) -> Once<A>
impl<A> Iterator for Once<A> { type Item = A; ... }
```

A variant: `std::iter::once_with(f)`.

The same, but **lazily**, through a closure.

```
fn once_with<A, F>(gen: F) -> OnceWith<F>
    where F: FnOnce() -> A
impl<A, F: FnOnce() -> A> Iterator for OnceWith<F> {
    type Item = A;
    ...
}
```

Note: generic with respect to the closure type `F`!

Call example:

```
let it = std::iter::once_with(|| 42);
println!("{}", it.next());
```

Closures

Three traits: FnOnce, FnMut, Fn

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

```
trait FnOnce<Args> {
    type Output;
    fn call_once(self, args: Args) -> Self::Output;
}

trait FnMut<Args>: FnOnce<Args> { // Sub-trait
    fn call_mut(&mut self, args: Args) -> Self::Output;
}

trait Fn<Args>: FnMut<Args> { // Sub-trait
    fn call(&self, args: Args) -> Self::Output;
}
```

- Rust automatically generates as many instances as possible when using the `| ... | ...` syntax.
- Captured variables are moved/mutably borrowed only if necessary.

Examples of APIs with closures

Iterator combinators

There exists many **combinators for iterators**.

Examples:

```
trait Iterator {  
    ...  
    fn map<B, F>(self, f: F) -> Map<Self, F>  
        where F: FnMut(Self::Item) -> B { ... }  
    fn filter<P>(self, predicate: P) -> Filter<Self, P>  
        where P: FnMut(&Self::Item) -> bool { ... }  
    ...  
}  
  
impl<B, I: Iterator, F> Iterator for Map<I, F>  
where F: FnMut(I::Item) -> B,  
{ type Item = B; ... }  
  
impl<I: Iterator, P> Iterator for Filter<I, P>  
where P: FnMut(&I::Item) -> bool,  
{ type Item = I::Item; ... }
```

We can write, for example:

```
(0..N).filter(|i| i % 2 == 0).map(|i| i * i).sum()
```

Zero-cost abstraction with iterators and closures

Let's see how the following piece of code gets compiled:

```
(0..N).filter(|i| *i % 2 == 0).map(|i| i * i + a).sum()
```

We first expand closures...

Zero-cost abstraction with iterators and closures

After expanding closures:

```
struct C1 { /* No captured variable */ }
impl FnOnce<(&i32,)> for C1 { type Output = i32; ... }
impl FnMut<(&i32,)> for C1 {
    fn call_mut(&mut self, args:(&i32,)) -> bool { *args.0 % 2 == 0 }
}

struct C2 { a: i32 }
impl FnOnce<(i32,>) for C2 { type Output = i32; ... }
impl FnMut<(i32,>) for C2 {
    fn call_mut(&mut self, args:(i32,)) -> i32 { args.0 * args.0 + self.a }
}

(0..N).filter(C1 {}).map(C2 { a: a }).sum()
```

We now inline `filter`, `map` and `sum`...

Zero-cost abstraction with iterators and closures

After inlining `filter`, `map` and `sum`:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 { a: a },
    iter:
        Filter {
            predicate: C1 {},
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    match it.next() {
        None => break,
        Some(i) => r += i
    }
}
```

We now inline `<Map>::next...`

Iterators

Closures

Trait objects

Zero-cost abstraction with iterators and closures

After inlining `<Map>::next`:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 { a: a },
    iter:
        Filter {
            predicate: C1 {},
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    let o = match it.iter.next() { None => None, Some(i) => Some(it.f.call_mut(i)) };
    match o {
        None => break,
        Some(i) => r += i
    }
}
```

Important: `call_mut` is a **direct call!**

We can merge the two `match` constructs, and inline `<C2>::call_mut...`

Iterators

Closures

Trait objects

Zero-cost abstraction with iterators and closures

After inlining `<Filter>::next` and simplifications:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 { a: a },
    iter:
        Filter {
            predicate: C1 {},
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    match it.iter.next() {
        None => break,
        Some(i) => r += i*i+it.f.a
    }
}
```

Similarly, we can inline `<Filter>::next` and `<C1>::call_mut...`

Zero-cost abstraction with iterators and closures

After inlining `<Filter>::next` and `<C1>::call_mut`:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 { a: a },
    iter:
        Filter {
            predicate: C1 {},
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    let o;
    loop { match it.iter.next() {
        None => { o = None; break },
        Some(i) => if i % 2 == 0 { o = Some(i); break }
    } }
    match o {
        None => break,
        Some(i) => r += i*i+it.f.a
    }
}
```

In the CFG graph, it is easy to see that the two `match` constructs can be merged...

Zero-cost abstraction with iterators and closures

We merged the two `match` constructs, and simplified:

```
/* Definition and instances of C1 and C2. */
...
let mut it = Map {
    f: C2 { a: a },
    iter:
        Filter {
            predicate: C1 {},
            iter: Range { start: 0, end: N }
        }
};
let mut r = 0;
loop {
    match it.iter.iter.next() {
        None => break,
        Some(i) => if i % 2 == 0 { r += i*i+it.f.a }
    }
}
```

We inline `<Range>::next()` and expands the parts of `it` that stay unchanged...

Zero-cost abstraction with iterators and closures

... and we finally obtain:

```
let mut range_start = 0;
loop {
    if range_start < N {
        let i = range_start;
        range_start += 1;
        if i % 2 == 0 { r += i*i+a }
    } else { break }
}
```

We started with this concise and “abstract” piece of code:

```
(0..N).filter(|i| i % 2 == 0).map(|i| i * i + a).sum()
```

And ended-up with a well-optimized version!

No optimization step was difficult, any optimizing compiler performs these!

An example of zero-cost abstraction!

Trait objects

Sometimes, a true **function type** is really needed.

- Because too much specialization increases generated code size.
- Because sometimes the set of required specialization is unbounded (ex. code generated after CPS transform).

Solution: some traits can be turned into a type: **trait object**.

- Can be seen as a bounded existential type.
- Syntax: `dyn Trait`
- Trait objects types are **unsized**.
 - Size cannot be determined statically, depends on the implementing type.
 - Only used with (fat) pointers: `Box<dyn Trait>`, `&dyn Trait`, ...
- Method calls are performed indirectly, through a vtable.
 - Some methods are incompatible: generics, using `Self` without a pointer...
⇒ only **object safe** traits are allowed in a trait object.

Trait objects

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators

Closures

Trait objects

What type would you use for continuations in CPS?

Trait objects

Rust: Generics
and Traits

Jacques-Henri
Jourdan

Iterators
Closures
Trait objects

What type would you use for continuations in CPS?

We could use `Box<dyn FnOnce(...) -> ...>`.

Conclusion: Rust uses the **same abstraction mechanism** for both static dispatch (i.e., templates in C++) and dynamic dispatch (i.e., virtual methods in C++).